

CS 4400

Computer Systems

LECTURE 3

Representing integers

Integer arithmetic

Encoding Integers

- Two different ways bits can be used to encode integers:
 - unsigned – only nonnegative numbers represented
 - signed – negative, zero, and positive values represented
- Both encodings represent a finite range of integers.

type declaration (C)	min	max
char	-128	127
unsigned char	0	255
short	-32768	32767
unsigned short	0	65535
int	-2147483648	2147483647
unsigned int	0	4294967295

Unsigned Integers

- Let vector $\vec{x} = [x_{w-1}, x_{w-2}, \dots, x_0]$ denote a w-bit integer value.
- Treat \vec{x} as a number written in binary notation to obtain the unsigned interpretation.

$$\text{B2U}_w(\vec{x}) = \sum_{i=0}^{w-1} x_i * 2^i$$

- $\text{UMin}_w = [00 \dots 00] = 0$
- $\text{UMax}_w = [11 \dots 11] = 2^w - 1$
- $\text{B2U}_w: \{0,1\}^w \rightarrow \{0, \dots, 2^w - 1\}$
- Bijection—associates a unique value to each w-bit vector.

Signed Integers

- The most common computer representation of signed integers is two's complement.

$$\text{B2T}_w(\vec{x}) = -x_{w-1} * 2^{w-1} + \sum_{i=0}^{w-2} x_i * 2^i$$

- Sign bit—the MSB, 1: negative and 0: nonnegative.
- $TMin_w = [10 \dots 00] = -2^{w-1}$
- $TMax_w = [01 \dots 11] = 2^{w-1} - 1$
- $\text{B2T}_w: \{0,1\}^w \rightarrow \{-2^{w-1}, \dots, 2^{w-1} - 1\}$
- Is B2Tw a bijection?

Exercise: Encoding Integers

- Let $w = 4$

Hex	Binary	$B2U_w$	$B2T_w$
0xA	[1010]	$2^3 + 2^1 = 10$	$-2^3 + 2^1 = -6$
0xB			
0xC			
0xD			
0xE			
0xF			

- Let $w = 8$. Compute $B2T_w$ for hex 0xAE.

Exercise: Encoding Integers

- Let $w = 4$

Hex	Binary	$B2U_w$	$B2T_w$
0xA	[1010]	$2^3 + 2^1 = 10$	$-2^3 + 2^1 = -6$
0xB	[1011]	$2^3 + 2^1 + 2^0 = 11$	$-2^3 + 2^1 + 2^0 = -5$
0xC	[1100]	$2^3 + 2^2 = 12$	$-2^3 + 2^2 = -4$
0xD	[1101]	$2^3 + 2^2 + 2^0 = 13$	$-2^3 + 2^2 + 2^0 = -3$
0xE	[1110]	$2^3 + 2^2 + 2^1 = 14$	$-2^3 + 2^2 + 2^1 = -2$
0xF	[1111]	$2^3 + 2^2 + 2^1 + 2^0 = 15$	$-2^3 + 2^2 + 2^1 + 2^0 = -1$

- Let $w = 8$. Compute $B2T_w$ for hex 0xAE.

More on Two's Complement

- The two's complement range is asymmetric.
- $UMax_w > 2 * TMax_w$. Why?
- Both encodings represent numeric value 0 the same way.
- The C standard does not require two's complement for signed integers.
 - Nearly all machines use it anyway. Does this affect portability?
 - See `limits.h` for constants delimiting ranges of different integer data types for a particular compiler and machine.
 - Other ways of representing signed integers?

Question

- When is $\sim x$ equivalent to each of the following expressions?
 - A. x
 - B. $-x$
 - C. $x + 1$
 - D. $-x + 1$
 - E. $-x - 1$

Signed-Unsigned Conversions

- Since both $B2U_w$ and $B2T_w$ are bijections, they have well-defined inverses, $U2B_w$ and $T2B_w$.
- Consider $U2T_w(\vec{x}) = B2T_w(U2B_w(\vec{x}))$
 - Takes number between 0 and 2^{w-1} , yields number between -2^{w-1} and $2^{w-1}-1$.
 - Both numbers have identical bit representations.
- Conversely, consider $T2U_w(\vec{x}) = B2U_w(T2B_w(\vec{x}))$
- Example (8-bit): $[10101010]$, 170 unsigned, -86 signed
- How do these functions affect signed and unsigned in C?

Unsigned and Signed in C

```
int x = -1;
unsigned ux = (unsigned) x; // ux is UMax_w
```

- In C, values are signed unless
 - explicitly type-cast (e.g., `(unsigned short)`)
 - put `U` in constant (e.g., `1234U`)
 - doesn't fit in `signed long`
 - starts `0x` and doesn't fit in `signed int`
- Use conversion codes `%d` (or `%i`), `%u` to print signed and unsigned decimal values, respectively.

Implicit Casts in C Expressions

- The values in a C expression are “promoted” to signed integer type before the expression is evaluated
 - Unsigned values smaller than int become signed
 - Unsigned values of int size (or larger) stay unsigned
 - Why?
- When signed and unsigned values of the same size are mixed in an expression, the signed values are converted into unsigned
- These language features interact poorly sometimes

Example: Unsigned and Signed

```
#include <stdio.h>
```

```
int main(void) {  
    int tx, ty;  
    unsigned ux, uy;
```

```
    tx = -96;  
    uy = 15;
```

```
    ux = (unsigned) tx; // explicit cast to unsigned  
    ty = uy;           // implicit cast to signed
```

```
    printf("%d, %d, %u, %u\n", tx, ty, ux, uy);
```

```
    printf("%d\n", ux + ty); // WHY = -81??
```

```
    return 0;
```

```
}
```

```
unix> gcc unsigned_signed.c
```

```
unix> ./a.out
```

```
-96, 15, 4294967200, 15
```

```
-81
```

unsigned_signed.c

Expanding Bit Representations

- A common operation is to convert between integers of different word sizes, retaining the same numeric value.
- To convert from smaller word size to larger:
 - for unsigned, simply add leading 0s – zero extension
 - for signed, add leading Xs such that X=MSB – sign extension
- Example:

```
short sx = 12345;    // 0x3039
short sy = -12345;   // 0xCFC7
int x = sx;          // 0x00003039
int y = sy;          // 0xFFFFCFC7
```

What does this program print?

```
#include <stdio.h>

int main (void)
{
    long a = -1;
    unsigned b = 1;
    printf ("%d\n", a > b);
    return 0;
}
```

Truncating Bit Representations

- To convert from larger word size to smaller (w -bit to k -bit, where $w > k$):
 - drop high-order $w-k$ bits – truncation
- Truncation of a number can alter its value, a form of overflow.

```
short x = (int) 12345; // 0x00003039, x is 12345  
short y = (int) 53191; // 0x0000CFC7, y is -12345
```

- For unsigned x , truncation to k -bit equivalent to $x \bmod 2^k$.
- For signed x ?

Advice on Unsigned

- Implicit casts are tricky (because they are easy to overlook) and can lead to bugs.
- To avoid such bugs, one might consider using only signed values.
 - Few languages other than C support unsigned values.
 - Java supports only signed values, requires two's complement, and guarantees that `>>` is an arithmetic shift.
- Unsigned values are very useful when thought of as a collection of bits (flags), with no math interpretation.

Unsigned Addition

- Consider w -bit unsigned values x and y , $0 \leq x, y \leq 2^w - 1$.
 - Representing the sum could require $w+1$ bits, $0 \leq x + y \leq 2^{w+1} - 2$
- In math, we cannot place any bound on the word size required to fully represent the results of arithmetic ops.
- Unsigned arithmetic is a form of modulo arithmetic.
 - Unsigned addition is equivalent to $(x + y) \bmod 2^w$.
 - `unsigned_add(x, y) = x + y`, if $x + y < 2^w$
 - `unsigned_add(x, y) = x + y - 2^w`, if $2^w \leq x + y < 2^{w+1}$
- Example: `unsigned short x = 65530 + 6; // x is 0`

Unsigned Overflow

- An arithmetic operation is said to *overflow* when the full integer result cannot fit within the limits of the data type.
- In C, overflow is not signaled as an error.
 - Some types of overflow may be signaled with a warning.
- We know that overflow has occurred during unsigned integer addition $s = x + y$, if $s < x$ (equivalently, if $s < y$).

- Example:

```
unsigned x = ~0;
unsigned y = 2;
unsigned s = x + y;
if(s < x) { ... } // overflow
```

Two's Complement Addition

- Consider w -bit values x and y , $-2^{w-1} \leq x, y \leq 2^{w-1}-1$.
 - Representing the sum could require $w+1$ bits, $-2^w \leq x + y \leq 2^w-2$
- We must truncate the result to w bits.
 - However, this is not as familiar as modulo arithmetic.
- The w -bit sum is the same as for unsigned addition.
$$\text{U2T}_w([(x + y) \bmod 2^w])$$
- Both positive and negative overflow can occur.
- *These overflows are undefined behavior in C and C++*

What does this program do?

```
#include <stdio.h>
#include <limits.h>

int foo (int x) {
    return (x+1) > x;
}

int main (void)
{
    printf ("%d\n", (INT_MAX+1) > INT_MAX);
    printf ("%d\n", foo(INT_MAX));
    return 0;
}
```

Output (depends on compiler version)

```
$ gcc overflow.c -O2
```

```
$ ./a.out
```

```
0
```

```
1
```

```
$ clang overflow.c -O2
```

```
$ ./a.out
```

```
0
```

```
1
```

Cases of Overflow

- *Negative overflow*—if $-2^w \leq x + y < -2^{w-1}$.
 - both x and y must be negative
 - a nonnegative integer is the result (counter to usual math)
 - `twoscomp_add(x, y) = x + y + 2w`
- *No overflow*—if $-2^{w-1} \leq x + y < 2^{w-1}$.
 - `twoscomp_add(x, y) = x + y`
- *Positive overflow*—if $2^{w-1} \leq x + y < 2^w$.
 - both x and y must be positive
 - a negative integer is the result (counter to usual math)
 - `twoscomp_add(x, y) = x + y - 2w`

Unsigned Multiplication

Consider w -bit unsigned values x and y , $0 \leq x, y \leq 2^w - 1$.

- The product could require $2w$ bits, $0 \leq x * y \leq (2^w - 1)^2$

We must truncate the result to w bits.

- In C, the low-order w bits are retained as the result.
- Equivalent to computing the product mod 2^w .

Example:

```
unsigned short x = 1 << 15; // 32768
x *= 3;                      // 32768
```

Two's Complement Multiplication

- Consider w -bit values x and y , $-2^{w-1} \leq x, y \leq 2^{w-1}-1$.
 - The product could require $2w$ bits, $-2^{2w-2} + 2^{w-1} \leq x * y \leq 2^{2w-2}$
- We must truncate the result to w bits.
 - However, this is not as familiar as for unsigned multiplication.
- The w -bit product is the same as for unsigned multiply.

$$U2T_w([(x * y) \bmod 2^w])$$

- *Signed multiplication overflows in C are undefined*

Multiplication by Powers of Two

- Integer multiplication used to be slow (≥ 12 cycles) compared to other integer operations.
 - addition, subtraction, bit-level ops, and shifts—1 cycle each
- An important (compiler) optimization was to replace multiplications by constant factors with shifts and adds.
- Let x be an integer. For any $k \geq 0$, $x * 2^k$ is equivalent to adding k 0's to the right of the bit representation of x .

- *Example:*

```
unsigned int = 11 << 3;           // 88
int = -11 << 3;                   // -88
```

Division by Powers of Two

- Integer division was also slow (≥ 30 cycles).
- Let x be an unsigned integer. For any $k \geq 0$, $x / 2^k$ is equivalent to adding k 0's to the left of the bit rep for x .
 - logical shift
- Let x be a signed integer. For any $k \geq 0$, $x / 2^k$ is equivalent to adding k b 's to the left of the bit rep for x .
 - b is the value of x 's MSB, arithmetic shift
 - What if $x < 0$?
- Example:

```
int x = 55 >> 3;    // 6
int y = -55 >> 3;   // -7 (should be -6)
```

Biasing

- If $x < 0$, integer division should round negative results up toward zero. Right shifting does not accomplish this.
- To correct for this improper rounding, we must “bias” the value before shifting.
 - First add $2^k - 1$ to x .
- For x , represented with two's complement and using arithmetic shifts, $x / 2^k$ is equivalent to

$$(x < 0 \text{ ? } (x + (1 \ll k) - 1) : x) \gg k$$

- Example:

```
int y = (-55 + (1 << 3) - 1) >> 3; // -6
```