

CS 4400

Computer Systems

LECTURE 5

Machine-level code

Accessing information

Arithmetic and logical operations

Machine-Level Rep of Programs

- High-level programming languages (PL) shield us from machine-level details (expressed in the assembly-code program).
 - exactly how memory is managed
 - functions are called
 - the low-level instructions used to carry out computation

Machine-Level Rep of Programs

- It is the job of a compiler to translate high-level programs to assembly code, so why must we understand it?
 - to analyze underlying inefficiencies in the code
 - to learn about the (hidden) run-time behavior of a program
- We will focus on how C programs are translated to assembly code using the IA32 (aka x86) instruction set.

Machine-Level Code

- Assembly code is very close to machine code.
 - but in a (more readable) text format, instead of binary
- Parts of processor state (hidden from high-level PL):
 - program counter (`%eip`), gives address of next instruction
 - integer registers, 8 named locations for 32-bit values
 - floating-point registers, 8 locations for floating-point data

Machine-Level Code

- A single machine instruction is very simple.
 - such as adding two numbers stored in registers, changing the program counter, etc.
- Compiler first job is to generate sequences of such instructions to implement high-level constructs (e.g., loops).

Machine-Level View of Memory

- Assembly code views memory simply as a large, byte-addressable array.
 - C arrays and structures are contiguous collections of bytes
 - no distinction between signed and unsigned data, pointers and integers, different types of pointers, ...

Machine-Level View of Memory

- Program memory contains:
 - object code for the program
 - (some info required by the OS)
 - a run-time stack for managing procedure calls and returns
 - blocks of memory allocated by the user (via `malloc`)

Example: Machine-Level Code

code.c

```
int accum = 0;

int sum(int x, int y) {
    int t = x + y;
    accum += t;
    return t;
}
```

code.s

```
...
sum:
    pushl %ebp
    movl %esp, %ebp
    movl 12(%ebp), %eax
    addl 8(%ebp), %eax
    addl %eax, accum
    movl %ebp, %esp
    popl %ebp
    ret
...
```

code.o

```
...55 89 e5 8b 45 0c 03
45 08 01 05 00 00 00 00
89 ec 5d c3...
```

- Assembly and object code from the text.
- gcc generates assembly code in “AT&T” format.

```
unix> gcc -S -O2 code.c
unix> ls
code.c
code.s
```

```
unix> gcc -c -O2 code.c
unix> ls
code.c
code.s
code.o
```


Data Formats

- An instruction has a suffix denoting the size of the operand.
(`movl`: move double word)
- Because Intel started as a 16-bit architecture:
 - “byte” 8-bit data type, instruction suffix is `b` (e.g., `char`)
 - “word” 16-bit data type, instruction suffix is `w` (e.g., `short`)
 - “double word” 32-bit data type, instruction suffix is `l` (e.g., `int`)
 - “quad word” 64-bit data type, instruction suffix is `q` (e.g., `long`)
- Suffix for single-precision floating point is `s`, for double precision suffix is `l`. Confusion with 4-byte integer?

32-Bit Registers

- 8 registers
- store integer data and pointers
- low-order bits may be independently read/written
- conventions for using “general purpose” registers to be covered

to access		double word	word	byte	byte
“general purpose”		%eax	%ax	%ah	%al
		%ecx	%cx	%ch	%cl
		%edx	%dx	%dh	%dl
		%ebx	%bx	%bh	%bl
		%esi	%si		
		%edi	%di		
stack pointer		%esp	%sp		
frame pointer		%ebp	%bp		

Operand Specifiers

- Most instructions have one or more operands.
 - specify source values to reference in performing operation
 - specify destination location into which to place the result
- Source values
 - immediate (constant), e.g., `$-577` or `$0x1F`
 - register, e.g., `%eax` (double-word op) or `%al` (byte op)
 - memory reference, e.g., `7(%eax)` (addr in `%eax+7`)
- Destination locations
 - register
 - memory reference

Addressing Modes

- $M[addr]$ denotes the value stored at byte address *addr*.
 - $R[reg-id]$ denotes the value of the contents of register *reg-id*.
-
- $0x2a3$, absolute, $M[0x2a3]$
 - $(\%eax)$, indirect, $M[R[\%eax]]$
 - $7(\%edx)$, base + displacement, $M[7 + R[\%edx]]$
 - $(\%eax, \%ecx)$, indexed, $M[R[\%eax] + R[\%ecx]]$
 - $7(\%eax, \%ecx)$, indexed, $M[7 + R[\%eax] + R[\%ecx]]$
 - $(, \%eax, 4)$, scaled indexed, $M[R[\%eax] * 4]$
 - $7(, \%eax, 4)$, scaled indexed, $M[7 + R[\%eax] * 4]$
 - $(\%eax, \%ecx, 4)$, scaled indexed, $M[R[\%eax] + R[\%ecx] * 4]$
 - $7(\%eax, \%ecx, 4)$, scaled indexed, $M[7 + R[\%eax] + R[\%ecx] * 4]$

Exercise: Addressing Modes

address	value
0x200	0x12
0x204	0x2a
0x208	0xd4
0x20c	0xfd

register	value
%eax	0x200
%ecx	0x41
%edx	0x4

operand	location	value
(%edx,%ecx,8)		
\$0x204		
%eax		
0x1f8(,%edx,4)		
(%eax)		
0x208		

Exercise: Addressing Modes

address	value
0x200	0x12
0x204	0x2a
0x208	0xd4
0x20c	0xfd

register	value
%eax	0x200
%ecx	0x41
%edx	0x4

operand	location	value
(%edx,%ecx,8)	0x2a3, absolute, M[0x2a3]	
\$0x204	(%eax), indirect, M[R[%eax]]	
%eax	7 (%edx), base + displacement, M[7 + R[%edx]]	
	(%eax,%ecx), indexed, M[R[%eax] + R[%ecx]]	
0x1f8(,%edx,4)	7 (%eax,%ecx), indexed, M[7 + R [%eax] + R[%ecx]]	
	(,%eax,4), scaled indexed, M[R[%eax] * 4]	
(%eax)	7 (,%eax,4), scaled indexed, M[7 + R[%eax] * 4]	
	(%eax,%ecx,4), scaled indexed, M[R[%eax] + R[%ecx] * 4]	
0x208	7 (%eax,%ecx,4), scaled indexed, M[7 + R[%eax] + R[%ecx] * 4]	

Exercise: Addressing Modes

address	value
0x200	0x12
0x204	0x2a
0x208	0xd4
0x20c	0xfd

register	value
%eax	0x200
%ecx	0x41
%edx	0x4

operand	location	value
(%edx,%ecx,8)	$M[0x4 + 0x41 * 8] = M[0x20c]$	0xfd (253)
\$0x204		
%eax		
0x1f8(,%edx,4)		
(%eax)		
0x208		

Exercise: Addressing Modes

address	value
0x200	0x12
0x204	0x2a
0x208	0xd4
0x20c	0xfd

register	value
%eax	0x200
%ecx	0x41
%edx	0x4

operand	location	value
(%edx,%ecx,8)	$M[0x4 + 0x41 * 8] = M[0x20c]$	0xfd (253)
\$0x204	—	0x204 (516)
%eax		
0x1f8(,%edx,4)		
(%eax)		
0x208		

Exercise: Addressing Modes

address	value
0x200	0x12
0x204	0x2a
0x208	0xd4
0x20c	0xfd

register	value
%eax	0x200
%ecx	0x41
%edx	0x4

operand	location	value
(%edx,%ecx,8)	$M[0x4+0x41*8]=M[0x20c]$	0xfd (253)
\$0x204	—	0x204 (516)
%eax	—	0x200 (512)
0x1f8(,%edx,4)		
(%eax)		
0x208		

Exercise: Addressing Modes

address	value
0x200	0x12
0x204	0x2a
0x208	0xd4
0x20c	0xfd

register	value
%eax	0x200
%ecx	0x41
%edx	0x4

operand	location	value
(%edx,%ecx,8)	$M[0x4+0x41*8]=M[0x20c]$	0xfd (253)
\$0x204	—	0x204 (516)
%eax	—	0x200 (512)
0x1f8(,%edx,4)	$M[0x1f8+0+0x4*4]=M[0x208]$	0xd4 (212)
(%eax)		
0x208		

Exercise: Addressing Modes

address	value
0x200	0x12
0x204	0x2a
0x208	0xd4
0x20c	0xfd

register	value
%eax	0x200
%ecx	0x41
%edx	0x4

operand	location	value
(%edx,%ecx,8)	$M[0x4+0x41*8]=M[0x20c]$	0xfd (253)
\$0x204	—	0x204 (516)
%eax	—	0x200 (512)
0x1f8(,%edx,4)	$M[0x1f8+0+0x4*4]=M[0x208]$	0xd4 (212)
(%eax)	$M[0x200]$	0x12 (18)
0x208		

Exercise: Addressing Modes

address	value
0x200	0x12
0x204	0x2a
0x208	0xd4
0x20c	0xfd

register	value
%eax	0x200
%ecx	0x41
%edx	0x4

operand	location	value
(%edx,%ecx,8)	$M[0x4 + 0x41 * 8] = M[0x20c]$	0xfd (253)
\$0x204	—	0x204 (516)
%eax	—	0x200 (512)
0x1f8(,%edx,4)	$M[0x1f8 + 0 + 0x4 * 4] = M[0x208]$	0xd4 (212)
(%eax)	$M[0x200]$	0x12 (18)
0x208	$M[0x208]$	0xd4 (212)

Data Movement Instructions

- Generality of the operand notation allows a simple move instruction to perform many different kinds of moves.
- `movl src, dest` (`movw` for 16-bit, `movb` for 8-bit)
 - `movl $0x4050, %eax ; immediate to register`
 - `movl %ebp, %esp ; register to register`
 - `movl $-17, (%esp) ; immediate to memory`
 - `movl %eax, -12(%ebp) ; register to memory`
- IA32 does not allow both operands to be memory locations.
 - How can we move the contents at address `0xa3` to address `0x7b`?

Question

Is this a valid IA32 instruction?

```
movb $0xF, (%b1)
```

Question

Is this a valid IA32 instruction?

```
movb $0xF, (%bl)
```

- Move Byte
 - From: Immediate value 0xf
 - To: M[%bl]

Question

Is this a valid IA32 instruction?

```
movw %ax, (%esp)
```


Question

Is this a valid IA32 instruction?

```
movw %ax, (%esp)
```

- Move Word (16-bit)
 - From: Word Register %ax
 - To: M[%esp]

Data Movement to/from Stack

- `pushl` and `popl` push data onto and pop data from the run-time stack. (more on the stack next week)
- `%esp` (the stack pointer) holds the address of the top stack element.
- `pushl %ebp` is equivalent to (why?):
`subl $4, %esp`
`movl %ebp, (%esp)`
- `popl %eax` is equivalent to (why?):
`movl (%esp), %eax`
`addl $4, %esp`

Exercise: Data Movement

```
int exchange(int *xp, int y)
{
    int x = *xp;
    *xp = y;
    return x;
}
```

```
movl 8(%ebp), %eax
movl 12(%ebp), %edx
movl (%eax), %ecx
movl %edx, (%eax)
movl %ecx, %eax
```

- What does the exchange function do?
- (Instructions for allocating / deallocating the stack frame are omitted.)
- When procedure begins, `xp` and `y` are stored at offsets 8 and 12 from the address in `%ebp` (the frame pointer).
- What do each of the assembly code instructions do?

Load Effective Address

- `leal` copies the address of the first operand to the destination (the second operand).
- Can be used to generate pointers.
- Can be used to compactly describe arithmetic operations.

Load Effective Address

- `leal 7(%edx,%edx,4),%eax`
 - $\%eax = 7 + \%edx + \%edx * 4$
 - if register `%edx` contains value x , then `%eax` is set to $5x + 7$

Load Effective Address

- Let `%eax` hold value x
- Let `%ecx` hold value y
- What value is stored in `%edx` after instruction
`leal 0xa(%eax,%ecx,4),%edx` ?

Load Effective Address

- Let `%eax` hold value x
- Let `%ecx` hold value y
- What value is stored in `%edx` after instruction
`leal 0xa(%eax,%ecx,4),%edx` ?

$$0xa + x + y * 4$$

Unary Operations

- One operand (register or memory location) serves as the source and destination.
 - `incl %eax`, increment value in register `%eax` by one
 - `decl (%eax)`, decrement value at `M[R[%eax]]` by one
 - `negl 6(%eax)`, negate value at `M[6 + R[%eax]]`
 - `notl (, %eax, 4)`, complement value at `M[R[%eax] * 4]`

Binary Operations

- Right operand (register or memory location) serves both as the first source and the destination.
 - Left operand (immed, reg, memory) is the second source.
 - Cannot have both operands as memory locations.
- `subl %eax, (%ecx)`
 - $M[R[\%ecx]] = M[R[\%ecx]] - R[\%eax]$

Binary Operations

- Other examples:
 - `addl` add
 - `imull` multiply
 - `xorl` exclusive or
 - `orl` or
 - `andl` and

Shift Operations

- First operand indicates by how much to shift.
 - immediate between 0 and 31 or in single-byte register `%cl`
- Second operand is value to shift.
- Arithmetic shifts
 - `sall $4, %ecx`, left shift equivalent to $R[\%ecx] = R[\%ecx] \ll 4$
 - `sarl`, right shift (fill with copies of MSB)
- Logical Shifts
 - `shll`, left shift, same as `sall`
 - `shrl`, right shift (fill with zeros)

Question

Which of the following is not equivalent to the C function?

```
int subtract(int x, int y) {  
    return y - x;  
}
```

- A. `movl 8(%ebp), %ecx`
`movl 12(%ebp), %eax`
`subl %ecx, %eax`
- B. `movl 8(%ebp), %eax`
`subl 12(%ebp), %eax`
- C. `movl 12(%ebp), %eax`
- D. `subl 8(%ebp), %eax`
- E. All are equivalent.
- F. More than one are not equivalent.

Exercise: Arithmetic

```
int arith(int x, int y, int z) {  
    int t1 = x + y;  
    int t2 = z * 48;  
    int t3 = t1 & 0xffff;  
    int t4 = t2 * t3;  
    return t4;  
}
```

```
movl 12(%ebp),%eax  
movl 16(%ebp),%edx  
addl 8(%ebp),%eax  
leal (%edx,%edx,2),%edx  
sall $4,%edx  
andl $65535,%eax  
imull %eax,%edx  
movl %edx,%eax
```

- The values of `x`, `y` and `z` are stored at offsets 8, 12 and 16 from the address in `%ebp`.
- Which assembly instruction(s) correspond to each C statement?

More Arithmetic Operations

- The two-operand `imull` generates a 32-bit product.
 - `imull src`, gives full 64-bit product of `%eax` and `src`, 2's comp
 - `mull src`, does the same for unsigned
 - both store the result in `%edx` (high order) and `%eax` (low order)

More Arithmetic Operations

- `cld`, sign extends the value in `%eax` to 64-bit
 - stores the result in `%edx` (high order) and `%eax` (low order)

More Arithmetic Operations

- `idivl src`, takes as a dividend the 64-bit value in `%edx` (high order) and `%eax` (low order), `src` is divisor
 - quotient stored in `%eax` and remainder in `%edx`
 - `divl` is the same except unsigned

Exercise: Data Movement

```
void decode1(int *xp, int *yp, int *zp) {  
    // FILL IN  
  
}
```

```
movl 8(%ebp),%edi  
movl 12(%ebp),%ebx  
movl 16(%ebp),%esi  
movl (%edi),%eax  
movl (%ebx),%edx  
movl (%esi),%ecx  
movl %eax, (%ebx)  
movl %edx, (%esi)  
movl %ecx, (%edi)
```

- The values of `xp`, `yp` and `zp` are stored at offsets 8, 12 and 16 from the address in `%ebp`.
- Fill in `decode1` such that it has an equivalent effect.

Exercise: Shift Operations

```
int shift_left2_rightn(int x, int n) {  
    x <<= 2;  
    x >>= n;  
    return x;  
}
```

```
movl 12(%ebp),%ecx ;move n to %ecx  
movl 8(%ebp),%eax  ;move x to %eax  
  
??                ;x<<=2  
  
??                ;x>>=n
```

- Fill in the above assembly code such that it is generated from the above C function.