

# CS 4400

## Computer Systems

---

### LECTURE 9

Structs and alignment

Buffer overflow

# Review: Structures

---

- In C, a user-defined type is accomplished with a `struct`.
- Example:

```
struct element {  
    char name[10];  
    char symbol[5];  
    float weight;  
    float mass;  
};
```

- Declaration of a structure variable

```
struct element e1;
```

- allocates contiguous storage for all structure members.
- at least  $10 + 5 + 2 * \text{sizeof(float)}$  bytes

# Review: Structures

---

- Use typedef to avoid the awkward two-word type.

```
typedef struct element {  
    char name[10];  
    char symbol[5];  
    float weight;  
    float mass;  
} ELT;
```

```
ELT e1;
```

- What is the difference in a structure and an array?

# Review: Structures

---

```
ELT e1;
```

```
ELT* elt_ptr = &e1;
```

- To access a member of the structure variable, use the dot `.` operator.

```
e1.mass = 3.0;
```

```
strcpy(e1.name, "hydrogen");
```

- As with objects in C++, the pointer operator `->` can be used with pointers to structures.

```
printf("%s", (*elt_ptr).symbol);
```

```
printf("%s", elt_ptr->symbol);
```

# Review: Structures

---

- A self-referential structure has a member that is a pointer of the same type as the structure itself.

```
typedef struct node {  
    int data;  
    struct node* next;  
} NODE;  
... x->next->next->data ...
```


# Structs

---

- The compiler maintains information about each structure.
  - indicating byte offset of each field

- Example (IA32):

```
struct rec {  
    int i;  
    int j;  
    int a[3];  
    int* p;  
};
```



- Generated code adds the appropriate offset.
  - suppose r (type struct rec \*) is in %edx, to copy element r->i to element r->j:  

```
movl (%edx),%eax  
movl %eax,4(%edx)
```

# Exercise: Structs

```
struct prob {  
    int* p;  
    struct {  
        int x;  
        int y;  
    } s;  
    struct prob* next;  
};
```

```
void sp_init(struct prob* sp) {  
    sp->s.x = _____ ;  
    sp->p = _____ ;  
    sp->next = _____ ;  
}
```

```
movl 8(%ebp),%eax  
movl 8(%eax),%edx  
movl %edx,4(%eax)  
leal 4(%eax),%edx  
movl %edx,(%eax)  
movl %eax,12(%eax)
```

- Offset of each field?
- Total number of bytes?
- Fill in function, given assembly code for its body.

# Question

---

What is the IA32 offset of field `f` in `struct d`?

- A. 0
- B. 4
- C. 8
- D. 12
- E. 16
- F. none of the above

```
struct a {  
    int b;  
    int c;  
};  
  
struct d {  
    struct a* e;  
    float f;  
};
```



# Unions

---

- Unions provide a way for a single object to be referenced according to multiple types.

- Example:

```
union u {  
    char c;  
    int i[2];  
    double v;  
} x;  
x.v = 4.5;  
printf("%d %d\n", x.i[0], x.i[1]);
```

- `sizeof(union u)` is the max size of any of its fields.
- Technically, you should only read the variant you wrote.

# Unions

```
unsigned f2u(float f) {  
    union {  
        float f;  
        unsigned u;  
    } temp;  
    temp.f = f;  
    return temp.u;  
}
```

```
movl 8(%ebp),%eax
```

- The byte offset of each field is 0.
- Example:

```
union rec {  
    char c;  
    int i[2];  
    double v;  
};    // 8 bytes
```
- Assembly code lacks any information about type.

# Alignment

---

- Some systems restrict the addresses allowed for primitive types—they must be a multiple of  $k$ .
- Alignment restrictions simplify the interface between processor and memory.
  - avoids a 4-byte `int` straddling two 4-byte memory blocks

# Alignment

---

- Linux/IA32 alignment convention:
  - addresses of 1-byte data types are not restricted
  - addresses of 2-byte data types are multiples of 2
  - addresses of larger data types are multiples of 4


# Struct Field Alignment

- The compiler may need to insert gaps in field allocation to ensure each structure element is aligned.


- Example:

```
struct S1 {  
    int i;  
    char c;  
    int j;  
};
```

- 9 bytes (unaligned):



- 12 bytes (aligned):



- Is a gap required if we make `char c` the third field?

# Exercise: Struct Alignment

---

- Given the Linux/IA32 alignment policy, how is each structure aligned?

```
struct P1 { int i; char c; int j; char d; };
```

```
struct P2 { int i; char c; char d; int j };
```

```
struct P3 { short w[3]; char c[3]; }
```

```
struct P4 { short w[3]; char* c[3]; }
```

```
struct P5 { struct P1 a[2]; struct P2 *p };
```

# Question

---

Given the Linux/IA32 alignment policy, what is the total number of bytes required for s?

- A. 12
- B. 16
- C. 20
- D. 24
- E. 28
- F. none of the above

```
struct {  
    char a[3];  
    short b;  
    double c;  
    char* d;  
} s;
```

# Question

---

If reordering of fields is allowed, is it possible to avoid padding at all in `s`?

```
struct {  
    char a[3];  
    short b;  
    double c;  
    char* d;  
} s;
```



# Packed Structs

---

- Many compilers support non-standard extensions for creating “packed” structs that contain no internal padding
- When and why is this useful? Harmful?

# Packed Structs

---

- For gcc: `#pragma pack(n)`
  - Struct fields will be aligned to the minimum of their natural alignment and n
  - So, `pack(1)` creates structs with no padding
  - Use `#pragma pack( )` to reset the compiler to normal padding behavior
  - Be careful: structs inside packed structs are not packed by default!

# Out-of-Bounds Memory References

---

- C does no bounds checking for array references.
  - Do any programming languages perform bounds checking?
- Recall that the run-time stack is used to store local variables, as well as, register values and return address.
- What happens when an out-of-bounds element of a local array is written?
  - program “state” is potentially corrupted
  - examples?

# Buffer Overflow

---

- A common source of state corruption.
- Typically: A char array is allocated to the stack, but a string is written which exceeds the allocated space.

```
char* gets(char* s) {  
    int c; char* dest = s;  
    while((c=getchar()) != '\n' && c != EOF)  
        *dest++ = c;  
    *dest = '\0';  
    if(c == EOF)  
        return NULL;  
    return s;  
}
```

```
void echo() {  
    char buf[4];  
    gets(buf);  
    puts(buf);  
}
```


Any potential  
problems with  
`gets`?

# Example: Buffer Overflow

```
void echo() {  
    char buf[4];  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    pushl %ebp           ;save to stack  
    movl %esp,%ebp       ;set new fr_ptr  
    subl $20,%esp        ;alloc space  
    pushl %ebx           ;save to stack  
    addl $-12,%esp       ;alloc more space  
    leal -4(%ebp),%ebx    ;buf is %ebp-4  
    pushl %ebx           ;push buf  
    call gets
```

- What values of buf will corrupt the saved value of %ebp?
- What values will corrupt the return address?
- Alternative string-input functions: fgets(), gets\_s(), getchar(), C++ I/O



# Exploit Code

---

- When the byte encoding of executable code is fed into a program as an input string, buffer overflow can be used to get a program do something it otherwise would not.
  - Also include extra bytes to overwrite the return address with the address of this exploit code.
  - The effect of ret is to jump to (and execute) the exploit code.

# Exploit Code

---

- In Lab 3, you will get first-hand experience mounting a buffer-overflow attack.
  - Requires deep understanding of run-time stack organization, byte ordering, and instruction encoding.

# Exercise: Buffer Overflow

```
char* getline() {  
    char buf[8];  
    char* result;  
    gets(buf);  
    result = malloc(strlen(buf)+1);  
    strcpy(result, buf);  
    return result;  
}
```

```
Disassembly of getline:  
push %ebp  
mov %esp, %ebp  
sub $16,%esp  
push %esi  
push %ebx  
add $0xffffffff4,%esp  
lea 0xffffffff8(%ebp),%ebx  
push %ebx  
... call gets ...
```

*return address*  
*saved %ebp*

...
08 04 86 43
bf ff fc 94 ← %ebp
...

- If input is 012345678901,
  - program terminates with seg-fault.
  - Error occurs during return of getline.
  - Fill in stack just before add, and then after call to gets.
- To where does the program try to return?
- What registers have corrupted values?