

CS 4400

Computer Systems

LECTURE 10

Capabilities and limitations of
compilers

Optimization blockers

Machine-independent optimizations

Optimization

- Writing efficient programs requires
 - selecting good data structures and algorithms
 - writing source code that the compiler can optimize
- Often there is a trade-off between readability and speed
 - one can program a simple insertion sort in minutes
 - a super-fast sorting routine can take days to code, debug
- When should program performance be traded for ease of implementation and maintenance?
- Optimizations are machine independent or dependent

Capabilities of Compilers

- By determining what values are computed and how they are used, optimizing compilers can often generate faster code than a compiler doing a straightforward translation
- Optimizing compilers exploit opportunities
 - to simplify expressions
 - to use a single computation in several places
 - to reduce the number of times a given computation is performed
- All of this must be done without changing the “observable behavior” of the program

Limitations of Compilers

- 1. The observable behavior must be maintained
- 2. An ahead-of-time compiler has very little knowledge about what will happen at run time
- 3. Optimizations must be performed quickly

```
void twiddle1(int* xp, int* yp) {  
    *xp += *yp;  
    *xp += *yp;  
}
```

```
void twiddle2(int* xp, int* yp) {  
    *xp += *yp * 2;  
}
```

Which function is more efficient?

Is the behavior of each identical?

Optimization Blocker: Aliasing

- An optimization blocker is a feature of the program's behavior that inhibits good code generation
- Memory aliasing is when a single memory location can be referenced with multiple identifiers
 - The compiler must assume that different pointers may designate the same place, or overlapping places, in memory.

```
void swap(int* xp, int* yp) {  
    *xp = *xp + *yp;    /* x+y */  
    *yp = *xp - *yp;    /* x+y-y=x */  
    *xp = *xp - *yp;    /* x+y-x=y */  
}
```

What if `xp` and `yp`
are equal?

Defeating Aliasing

- Make a copy of an aliased variable, and modify the copy
 - Only when this works correctly, of course
- Assert that pointers are not aliases
 - But be careful: pointers can refer to overlapping storage even if they aren't equal
- As a last resort: Tag pointer parameters with the “restrict” qualifier
 - Example: `alias2.c`

Optimization Blocker: Function Calls

```
int counter = 0;
```

```
int f(int x) { return counter + x; }
```

```
int func1(int x) {  
    return f(x) + f(x) + f(x);  
}
```

```
int func2(int x) { return 4 * x; }
```

How are `func1`
and `func2`
different?

- Function `f` has a *side effect*—modifying part of the global program state.
- Most compilers do not try to determine whether a function is free of side effects.
 - They simply assume the worst case.

```

typedef int data_t;                                /* change as needed for float, ... */

typedef struct {
    int len;
    data_t* data;
} vec_rec, * vec_ptr;                             /* typedefs struct, pointer to struct */

vec_ptr new_vec(int len) { /* create vector of specified length */
    vec_ptr result = (vec_ptr) malloc(sizeof(vec_rec));
    if(!result) return NULL; /* cannot allocate storage */
    result->len = len;
    if(len > 0) {
        data_t* data = (data_t*) calloc(len, sizeof(data_t));
        if(!data) { /* cannot allocate storage */
            free((void*) result);
            return NULL;
        }
        result->data = data;
    }
    else result->data = NULL;
    return result;
}

/* retrieve vector element and store at dest */
int get_vec_element(vec_ptr v, int index, data_t* dest) {
    if(index < 0 || index >= v->len) /* bounds checking */
        return 0;
    *dest = v->data[index];
    return 1;
}

int vec_length(vec_ptr v) { return v->len; };

```


Example: Vector ADT

```
#define IDENT 0      /* 0,+ sums elements of vector */
#define OPER +      /* change to 1,* for product */

void combine1(vec_ptr v, data_t* dest) {
    int i;

    *dest = IDENT;
    for(i = 0; i < vec_length(v); i++) {
        data_t val;
        get_vec_element(v, i, &val);
        *dest = *dest OPER val;
    }
}
```


Performance Measurements

- Among compilers, gcc is considered to be pretty good at optimizing, though not the best
- *Unoptimized*—code suitable for stepping through with debugger, closely matches source code.
 - `-O1` enables basic optimizations
- *CPE* measures the number of clock cycles per element.
 - appropriate for programs that perform a repetitive computation
 - (e.g., processing pixels, computing elts of matrix product)
 - not necessarily cycles per iteration

Loop Inefficiency

- Observe that `combine1` calls `vec_length` as the test condition on every iteration of the loop.
- However, the vector length does not change.
 - As we know, the compiler will not move the function call.
 - The programmer must explicitly perform this optimization.
- *Code motion* optimization:
 - Identify a computation that is performed repeatedly, but whose result does not change.
 - Move the computation so that it does not get executed as often.

Example: Loop Inefficiency

```
/* move call to vec_length out of loop */
void combine2(vec_ptr v, data_t* dest) {
    int i;
    int length = vec_length(v);

    *dest = IDENT;
    for(i = 0; i < length; i++) {
        data_t val;
        get_vec_element(v, i, &val);
        *dest = *dest OPER val;
    }
}
```


Question

We express relative performance as a ratio of the form:

$$\frac{T_{old} = \text{time of the original version}}{T_{new} = \text{time of the modified version}}$$

Which of the following are true?

- A. A ratio of 0 means no improvement, 1 means slight improvement, 2 means significant improvement.
- B. The ratio will never be less than 1.
- C. The CPEs is 12.00 for `combine1` and 8.03 for `combine2`. Thus, the performance ratio is about 1.5.

Question

What is the total number of function calls in this loop? (Assume the x is 10 and y is 100.)

```
for(i = min(x, y); i < max(x, y); incr(&i, 1))  
    t += square(i);
```

- A. 4
- B. between 50 and 100
- C. between 101 and 200
- D. more than 200

Reducing Procedure Calls

- Procedure calls incur overhead and block optimizations.
- `get_vec_element` is called on every loop iteration.
 - especially costly procedure call because of bounds checking
 - simple analysis shows all array references to be valid

```
data_t* get_vec_start(vec_ptr v) { return v->data; }

/* direct access to vector data */
void combine3(vec_ptr v, data_t* dest) {
    int i;
    int length = vec_length(v);
    data_t* data = get_vec_start(v);

    *dest = IDENT;
    for(i = 0; i < length; i++)
        *dest = *dest OPER data[i];
}
```

Reducing Procedure Calls

- How does this transformation affect the modularity?
- The CPE improvement is up to a factor of 1.3X.
 - ratio $T_{old} / T_{new} = 8.03 / 6.01 = 1.34$
 - what is the factor if there is no improvement?
- Modest improvement, but call is bottleneck for future opts.
- Compromise modularity and abstraction for speed, but only if performance is a significant issue.

Reducing Memory References

```
# (x86-64 floating-pt code)
# combine3, data_t is float, OPER is *
# dest in %rbp, data in %rax, i in %rdx, length in %r12
.L498:
    movss (%rbp), %xmm0           # read *dest
    mulss (%rax,%rdx,4), %xmm0    # multiply by data[i]
    movss %xmm0, (%rbp)          # write *dest
    addq $1, %rdx                # i++
    cmpq %rdx,%r12               # compare i to length
    jg .L498                     # if i<length, goto loop
```

- The value being computed is accumulated in the location designated by pointer `dest`, memory read/write required.
- Possible to avoid so many reads and writes of memory?
 - value written is read on next iteration

Reducing Memory References

```
/* accumulate result in local variable */  
void combine4(vec_ptr v, data_t* dest) {  
    int i;  
    int length = vec_length(v);  
    data_t* data = get_vec_start(v);  
    data_t acc = IDENT;  
  
    for(i = 0; i < length; i++)  
        acc = acc OPER data[i];  
    *dest = acc;  
}
```

(AKA “scalar replacement”)

```
# combine4  
# data in %rax,  
# acc in %xmm0,  
# i in %rdx,  
# length in %rbp  
.L488:  
    mulss (%rax,%rdx,4),%xmm0  
    addq $1, %rdx  
    cmpq %rdx,%rdp  
    jg .L488
```


Will Compiler Reduce Refs?

- Is scalar replacement an optimization the compiler will perform automatically?
 - not in this case, because of potential memory aliasing
- Consider vector $v = [2, 3, 5]$, OPER is $*$, and calls
 - `combine3(v, get_vec_start(v)+2)` results in `[2, 3, 36]`
 - `combine4(v, get_vec_start(v)+2)` results in `[2, 3, 30]`
- An optimizing compiler cannot make a judgment about the conditions under which a function might be used. Thus, it is obliged to preserve its exact functionality.

Loop Unrolling

- Some loops have such a small body that most of the execution time is spent updating the loop-counter variable and testing the loop-exit condition.
- It is more efficient to unroll such loops, putting two or more copies of the loop body in a row.
- Then, avoid setting and testing the loop counter in every loop body, reducing “loop overhead”.
- How should the new loop update/exit compare to original?

Example: Loop Unrolling

BEFORE

```
L1:  x ← M[i]
      s ← s + x
      i ← i + 4
      if i < n goto L1
L2:
```

AFTER

```
L1: x ← M[i]
      s ← s + x
      x ← M[i+4]
      s ← s + x
      i ← i + 8
      if i < n goto L1
L2: goto L2
L2:
```

- Will this work if the original loop iterated an odd number of times?
- How can we accommodate an odd number of iterations?
- How can we modify our strategy to unroll by a factor of K?
- Will the optimizing compiler perform loop unrolling automatically?

Exercise: Loop Unrolling

```
void inner_prod(vec_ptr u, vec_ptr v, data_t *dest) {
    int i;
    int length = vec_length(u);
    data_t *udata = get_vec_start(u);
    data_t *vdata = get_vec_start(v);
    data_t sum = (data_t) 0;

    for (i = 0; i < length; i++) {
        sum += udata[i] * vdata[i];
    }

    *dest = sum;
}
```

Perform 4-way loop unrolling

Summary

- To effectively use optimizing compilers, programmers must know the capabilities and limitations.
- Machine-independent optimizations:
 - code motion
 - reducing procedure calls
 - reducing memory references
 - loop unrolling (its machine dependence to be revisited)
- The programmer does have to help the optimizing compiler by dealing with optimization blockers.