

# CS 4400

## Computer Systems

---

### LECTURE 12

The memory hierarchy

Locality

Cache memory

# Memory Model

---

- Up to this point, our model of the memory system has been overly simple
  - Memory: linear array
  - Access: constant time


# Memory Model

---

- Memory system is a ***hierarchy*** of storage devices with different capacities, costs, and access times.
- Why do we care?
  - Data movement is expensive
- Our programs must have good ***locality***.

# The Memory Hierarchy

---



# Storage Technologies

---

- Two varieties of random-access memory (RAM)
- Static: used for on-and off-chip caches
  - fast, expensive, typically a few MB
  - Lower densities than DRAM cells
    - hence more expensive and consume more power
- Dynamic: used for main memory
  - Slower, but larger
  - Lose charge within 10-100 ms / memory must be refreshed
  - Organized as 2D arrays

# Storage Technologies

---

- DRAMs and SRAMs are ***volatile***—they lose their information if the supply voltage is turned off.
- Nonvolatile memories are called read-only memory (ROM), even though some types can be written to
- Disks large but reading/writing slow
  - Disk — size: ~1-10 TB, access: milliseconds
  - SSD — size: ~128-1024 GB, access: microseconds
  - RAM — size: ~8-128 GB, access: nanoseconds

# Locality

---

- Well-written programs tend to access items that are “near” other recently-accessed items.
  - This principle of **locality** has enormous impact on the design and performance of hardware and software systems.
  - A program has good locality of reference if it can reuse data while in the upper levels of memory.

# Locality

---

- ***Temporal locality***—accessing recently-referenced data
- ***Spatial locality***—accessing data with memory addresses near those of recently-referenced data



# Exercise: Locality

---

```
int sumvec(int v[N]) {  
    int i, sum = 0;  
  
    for(i = 0; i < N; i++)  
        sum += v[i];  
  
    return sum;  
}
```

- Does the reference to `sum` have good locality?
  - If so, what kind(s)?
- Does the reference to `v` have good locality?
  - If so, what kind(s)?

# Exercise: Locality

---

```
int sumarrayrows(int a[M][N]) {  
    int i, j, sum = 0;  
  
    for(j = 0; j < N; j++)  
        for(i = 0; i < M; i++)  
            sum += a[i][j];  
  
    return sum;  
}
```

- Does the reference to `a` have good locality?
  - If so, what kind(s)?
- What if we interchange the `i`- and `j`-loops?

# Cache Memory

---

- Caches are small fast memories that hold blocks of the most recently accessed instructions and data.

# Cache Memory

---

- When the processor requires the datum at memory address  $x$ , it first looks in the cache.
- If  $x$  is in the cache, a **cache hit** occurs.
- If  $x$  is not in the cache, a **cache miss** occurs. The processor fetches  $x$  from main memory, placing a copy of  $x$  in the cache.
- Placing  $x$  in the cache may mean displacing (**evicting**) another datum from the cache.

# Simple Cache

## (fully associative cache)

---

- A **fully associative** cache contains set of storage locations that can hold any data element
- Cached data is stored in **blocks** or **cache lines**
  - Every cache has a fixed block size  **$B$**
- A cache of size  **$C$** , contains  **$E = C/B$**  cache lines
  - Each cache line stores a base address and the block of data

# Eviction Policy

---

- When a cache is full and we need to make space, we evict another data element using an ***eviction policy*** or ***replacement policy***
- Common policies:
  - ***Random***: Choose randomly which block to replace.
  - ***First-in first-out (FIFO)***: Choose to replace the block that has resided in the cache set the longest.
  - ***Least-recently used (LRU)***: Choose to replace the block that has been unused the longest.
- Is there an optimal policy?

# Basics of Associative Sets (direct-mapped cache)

---

- A ***direct-mapped*** cache contains the same  $E = C/B$  cache lines as a fully associative
- However, now every address maps exclusively to a single cache line
- Advantages/Disadvantages?
- Example

# Cache Associativity

---


direct-mapped

$E=1, S=8$




2-way associative

$E=2, S=4$



4-way associative

$E=4, S=2$



fully associative


$E=8, S=1$





# Cache Organization

CACHE:



blocksize  $B$


*a cache block*

- The binary expansion of a memory address is divided into a memory block address and an offset within the block.
- The cache blocksize  $B$  indicates how many contiguous bytes of memory are copied to cache.


MEMORY ADDRESS:



# Cache Organization




MEMORY ADDRESS:




- A memory address is further divided into a cache set and a tag.
- The number of cache sets  $S$  indicates how many sets are in the entire cache.

# Cache Organization



MEMORY ADDRESS:



- The capacity of the entire cache is  $C = E * B * S$ .
- Why is the organization of the cache so seemingly complex?

# Mapping Memory to Cache

---

- A cache contains  $C/B$  *cache lines*, and each line may be empty or may be occupied by a memory block.
- No two cache lines may contain the same memory block. Why?
- A *cache set* is a collection of  $E$  lines that a particular memory block may occupy in cache.
- The *tag* disambiguates memory blocks within the same cache set.

# Cache-Replacement Policy

---

- On a cache miss, the ***replacement policy*** selects which frame in a cache set to update with the new memory block.

# Cache-Miss Categories

---

- A ***compulsory (or cold) miss*** occurs on the very first access to a memory block.
  - Why?
  - Possible to avoid?

# Cache-Miss Categories

---

- A ***capacity miss*** occurs when accessing a block that previously resided in cache, but was replaced because the cache cannot hold all of the data needed to execute a program.
  - Possible to avoid?

# Cache-Miss Categories

---

- A ***conflict miss*** occurs when accessing a block that previously resided in cache, but was replaced because too many blocks map to the same cache set.
  - Possible to avoid?



# Cache-Miss Categories

---

- A capacity miss in an LRU set-associative cache with capacity  $C$  is also a miss in an LRU fully-associative cache with capacity  $C$ .
- A conflict miss in an LRU set-associative cache with capacity  $C$  is a hit in an LRU fully-associative cache with capacity  $C$ .
- *Why wouldn't we make all caches fully-associative, to avoid conflict misses?*

# Question


---

Consider a direct-mapped cache with 512 total bytes and 16-byte blocks. How many cache sets are there?

- A. 8 sets
- B. 16 sets
- C. 32 sets
- D. 64 sets
- E. There is not information to determine the number of sets.
- F. None of the above

# Exercise: Cache Parameters

- **C** – cache capacity (total number of bytes)
  - **B** – cache block size
  - **E** – cache associativity (number of blocks per set)
  - **S** – number of cache sets
  - **m** – number of main memory address bits
  - **t** – number of tag bits ( $m - s - b$ )
  - **s** – number of set index bits ( $\log_2 S$ )
  - **b** – number of block offset bits ( $\log_2 B$ )
- 
- What are **S**, **t**, **s**, and **b**?
    - $m = 32$ , **C** = 1024, **B** = 4, **E** = 1
    - $m = 32$ , **C** = 1024, **B** = 8, **E** = 4
    - $m = 32$ , **C** = 1024, **B** = 32, **E** = 32



# Question

---

Consider a 4-way set-associative cache with 512 total bytes and 16-byte blocks. Which of the following addresses of memory blocks map to cache set 5?

- A. 0x0836
- B. 0x0845
- C. 0x0877
- D. 0x0900
- E. Two or more of the above
- F. None of the above

# Example: DM Cache

```
float dotprod(float x[8], float y[8]) {  
    float sum = 0.0;  
    int i;  
  
    for(i = 0; i < 8; i++)  
        sum += x[i] + y[i];  
  
    return sum;  
}
```

“Thrashing” occurs—the cache is repeatedly loading and evicting the same sets of blocks.

How can thrashing be prevented?

- $E=1, B=16, C=32, S=2$
- `sizeof(float)` is 4, `x` starts at address 0, `y` starts at address 32

element	addr	set	miss?	element	addr	set	miss?
x[0]	0	0	cold	y[0]	32	0	cold
x[1]	4	0	conflict	y[1]	36	0	conflict
x[2]	8	0	conflict	y[2]	40	0	conflict
x[3]	12	0	conflict	y[3]	44	0	conflict
x[4]	16	1	cold	y[4]	48	1	cold
x[5]	20	1	conflict	y[5]	52	1	conflict
x[6]	24	1	conflict	y[6]	56	1	conflict
x[7]	28	1	conflict	y[7]	60	1	conflict

# Exercise: Associative Cache

- $E=2, B=4, C=64, S=8, m=13$
- Accesses are to 1-byte words
- Which bits (12 to 0) are used to determine:
  - cache block offset?
  - cache set index?
  - cache tag?
- Suppose address  $0x0E34$  is referenced. Hit or miss?
  - Reference to  $0x0DD5$ ?
  - Reference to  $0x1FE4$ ?

LINE (or FRAME) 1

set	tag	valid	byte0	byte1	byte2	byte3
0	09	1	86	30	3F	10
1	45	1	60	4F	E0	23
2	EB	0	--	--	--	--
3	06	0	--	--	--	--
4	C7	1	06	78	07	C5
5	71	1	0B	DE	18	4B
6	45	1	A0	B7	26	2D
7	46	0	--	--	--	--

LINE (or FRAME) 2

set	tag	valid	byte0	byte1	byte2	byte3
0	00	0	--	--	--	--
1	38	1	00	BC	0B	37
2	0B	0	--	--	--	--
3	32	1	12	08	7B	AD
4	05	1	40	67	C2	3B
5	6E	0	--	--	--	--
6	F0	0	--	--	--	--
7	DE	1	12	C0	88	37