

CS 4400

Computer Systems

LECTURE 15

Static libraries

Relocation

Shared libraries and dynamic linking

Looking at Object File Symbol Tables

- Use the “nm” command
- Example

Static Libraries

- ***Static library***: A collection of .o files glued together
 - linker copies only the object modules that the program refs
 - C example: defs of `printf`, `strcpy`, and `rand` are in `libc.a`

```
> gcc main.c /usr/libm.a /usr/libc.a
```

Static Libraries

- Why not put all library functions in a single module?
- Why not put each library function in its own module?

Resolving References

- The linker scans relocatable objects *left to right* as they appear on the command line.
 - Compiler driver automatically translates any `.c` files to `.o` files

Resolving References

- During the scan, the linker maintains:
 - **E**: a set of relocatable objects to be merged
 - **U**: a set of unresolved symbols (referred to but not yet defined)
 - **D**, a set of symbols that have been defined previously
 - initially sets **E**, **U**, and **D** are empty

Scanning Input Object Files

- For each input object file or library file \mathbb{f}
 - if \mathbb{f} is an object file
 - add \mathbb{f} to \mathbb{E}
 - update \mathbb{U} and \mathbb{D} to reflect the symbols and references in \mathbb{f}

\mathbb{E} = {relocatable objects}
 \mathbb{U} = {unresolved symbols}
 \mathbb{D} = {symbols defined}

Scanning Input Object Files

- For each input object file or library file \mathbb{f}
 - if \mathbb{f} is an object file
 - add \mathbb{f} to \mathbb{E}
 - update \mathbb{U} and \mathbb{D} to reflect the symbols and references in \mathbb{f}
 - if \mathbb{f} is a library
 - if member m defines a symbol in \mathbb{U}
 - add m to \mathbb{E}
 - update \mathbb{U} and \mathbb{D} to reflect defs and refs in m
 - iterate over all members until \mathbb{U} and \mathbb{D} no longer change
 - Discard any objects from \mathbb{f} not contained in \mathbb{E}

\mathbb{E} = {relocatable objects}
 \mathbb{U} = {unresolved symbols}
 \mathbb{D} = {symbols defined}

Scanning Input Object Files

- For each input object file or library file f
 - if f is an object file
 - add f to E
 - update U and D to reflect the symbols and references in f
 - if f is a library
 - if member m defines a symbol in U
 - add m to E
 - update U and D to reflect defs and refs in m
 - iterate over all members until U and D no longer change
 - Then discard any objects from f not contained in E
- If U is nonempty when linker finishes scanning,
ERROR

$E = \{\text{relocatable objects}\}$ $U = \{\text{unresolved symbols}\}$ $D = \{\text{symbols defined}\}$

Example: Scanning Input Files

```
unix> gcc ./libvector.a main2.c  
/tmp/cc9XH6Rp.o: In function `main':  
/tmp/cc9XH6Rp.o(.text+0x18): undefined reference to `addvec'
```

- If the library (which defines a symbol) appears on the command line before the object file (which references the symbol), the reference cannot be resolved.

Example: Scanning Input Files

- Libraries can be repeated on the command line as needed to satisfy dependencies.
 - Suppose that `foo.c` calls a function in `libx.a` that calls a function in `liby.a` that calls a function in `libx.a`.

```
unix> gcc foo.c libx.a liby.a libx.a
```

Exercise: Scanning Input Files

- Let $a \rightarrow b$ denote that a depends on b
 - i.e., b defines a symbol that is referenced by a
- Give the minimal command line that will allow the static linker to resolve all symbol references.

$p.o \rightarrow libx.a$

$p.o \rightarrow libx.a \rightarrow liby.a$

$p.o \rightarrow libx.a \rightarrow liby.a$
and $liby.a \rightarrow libx.a \rightarrow p.o$

Exercise: Scanning Input Files

- Let $a \rightarrow b$ denote that a depends on b
 - i.e., b defines a symbol that is referenced by a
- Give the minimal command line that will allow the static linker to resolve all symbol references.

$p.o \rightarrow libx.a$ $> gcc\ p.o\ libx.a$

$p.o \rightarrow libx.a \rightarrow liby.a$ $> gcc\ p.o\ libx.a\ liby.a$

$p.o \rightarrow libx.a \rightarrow liby.a$ $> gcc\ p.o\ libx.a$
and $liby.a \rightarrow libx.a \rightarrow p.o$ $liby.a\ libx.a$

Relocation

- The linker merges the input modules and assigns run-time addresses to each symbol

Relocation

- Step 1: relocate sections and symbol definitions
 - merge all sections of the same type into a new aggregate section
 - assign run-time addresses to new aggregate sections
 - assign run-time addresses to each symbol defined

Relocation

- Step 2: relocate symbol references within sections
 - modify every symbol reference in bodies of the code and data sections so that they point to the correct run-time addresses
 - linker relies on relocation entries `.rel.text` and `.rel.data` to perform this step

Relocating Symbol References

```
typedef struct {
    int offset;      /* offset of ref to relocate */
    int symbol:24,   /* symbol ref should point to */
        type:8;     /* relocation type */
} Elf32_Rel;
```

*format of ELF
relocation
entry*

relocation algorithm

```
foreach section s
    foreach relocation entry r {
        refptr = s + r.offset; /* ptr to ref to be relocated */

        if(r.type == R_386_32) /* relocate an absolute addr */
            *refptr = (unsigned) (ADDR(r.symbol) + *refptr);

        if(r.type == R_386_PC32) { /* relocate a PC-relative ref */
            refaddr = ADDR(s) + r.offset; /* ref's run-time addr */
            *refptr = (unsigned) (ADDR(r.symbol) + *refptr - refaddr);
        }
    }
}
```

Assumptions: `s` is an array of bytes, `r` has type `Elf32_Rel`, and linker has already chosen run-time addresses for each section (`ADDR(s)`) and each symbol (`ADDR(r.symbol)`)

Relocating Absolute References

```
int* bufp0 = &buf[0];
```

- `bufp0` will be:
 - stored in `.data` of `swap.o`
 - initialized to the address of a global array
- Thus, the value of `bufp0` must be relocated

Relocating Absolute References

```
int* bufp0 = &buf[0];
```

- `bufp0` will be:
 - stored in `.data` of `swap.o`
 - initialized to the address of a global array
- Thus, the value of `bufp0` must be relocated

```
00000000 <bufp0>:  
0: 00 00 00 00    int* bufp0 = &buf[0];  
0: R_386_32 buf    relocation entry
```

*Disassembled
listing of the
.data section
(from swap.o)*

```
r.offset = 0  
r.symbol = buf  
r.type = R_386_32
```

Relocating Absolute References

```
r.offset = 0, r.symbol = buf, r.type = R_386_32
```

- Assume:

```
ADDR(buf) = 0x8049454
```

Relocating Absolute References

```
r.offset = 0, r.symbol = buf, r.type = R_386_32
```

- Assume:

```
ADDR(buf) = 0x8049454
```

- Linker updates the reference:

```
*refptr = (unsigned)(ADDR(r.symbol) + *refptr)  
          = (unsigned)(0x8049454 + 0) = 0x8049454
```

Relocating Absolute References

```
r.offset = 0, r.symbol = buf, r.type = R_386_32
```

- Assume:

```
ADDR(buf) = 0x8049454
```

- Linker updates the reference:

```
*refptr = (unsigned)(ADDR(r.symbol) + *refptr)  
          = (unsigned)(0x8049454 + 0) = 0x8049454
```

- Linker decides that at run time `bufp0` will be located at `0x804945c` and will be initialized to `0x8049454`, the run-time address of the `buf` array.

```
0804945c <bufp0>:  
804945c: 54 94 04 08
```

disassembled .data listing (from executable object file)

Relocating PC-Relative References

opcode

reference (-4) biased because PC always points to next instruction

```
6: e8 fc ff ff ff  call 7 <main+0x7>  swap();  
7: R_386_PC32 swap                                relocation entry
```

disassembled call instruction (from main.o)

```
r.offset = 7  
r.symbol = swap  
r.type = R_386_PC32
```

Relocating PC-Relative References

```
r.offset = 7, r.symbol = swap, r.type = R_386_PC32
```

- Assume:

```
ADDR(.text) = 0x80483b4
```

```
ADDR(swap) = 0x80483c8
```


Relocating PC-Relative References

```
r.offset = 7, r.symbol = swap, r.type = R_386_PC32
```

- Assume:

```
ADDR(.text) = 0x80483b4
```

```
ADDR(swap) = 0x80483c8
```

- First, linker computes run-time address of the reference:

```
refaddr = ADDR(s) + r.offset  
         = 0x80483b4 + 0x7  
         = 0x80483bb
```

Relocating PC-Relative References

```
r.offset = 7, r.symbol = swap, r.type = R_386_PC32
```

- Assume:

```
ADDR(.text) = 0x80483b4
```

```
ADDR(swap) = 0x80483c8
```

- First, linker computes run-time address of the reference:

```
refaddr = ADDR(s) + r.offset
```

```
          = 0x80483b4 + 0x7
```

```
          = 0x80483bb
```

- Then, linker updates the reference from its current value (-4) so that it will point to the swap routine at run time:

```
*refptr = (unsigned)(ADDR(r.symbol) + *refptr - refaddr)
```

```
          = (unsigned)(0x80483c8 + (-4) - 0x80483bb)
```

```
          = 0x9
```

Relocating PC-Relative References

```
r.offset = 7, r.symbol = swap, r.type = R_386_PC32
```

- Assume:

```
ADDR(.text) = 0x80483b4
```

```
ADDR(swap) = 0x80483c8
```

- First, linker computes run-time address of the reference:

```
refaddr = ADDR(s) + r.offset
```

```
          = 0x80483b4 + 0x7
```

```
          = 0x80483bb
```

- Then, linker updates the reference from its current value (-4) so that it will point to the swap routine at run time:

```
*refptr = (unsigned)(ADDR(r.symbol) + *refptr - refaddr)
```

```
          = (unsigned)(0x80483c8 + (-4) - 0x80483bb)
```

```
          = 0x9
```

```
80483ba: e8 09 00 00 00 call 80483c8 <swap>
```

ELF Executable Object File Format

ELF header	describes overall format, includes entry point (addr of 1st instruction)
Segment header table	maps contiguous file sections to run-time memory sections
<code>.init</code>	describes function <code>_init</code> , to be called by program's initialization code
<code>.text</code>	read-only memory segment (code segment)
<code>.rodata</code>	
<code>.data</code>	
<code>.bss</code>	read/write memory segment (data segment)
<code>.symtab</code>	symbol table, debug info not loaded into memory
<code>.debug</code>	
<code>.line</code>	
<code>.strtab</code>	
section header table	**Notice the lack of <code>.rel.text</code> and <code>rel.data</code> sections. Why?

Loading Executable Object Files

\$./p

- Because `p` is not a built-in shell command, the shell assumes that `p` is an executable object file.

Loading Executable Object Files

\$./p

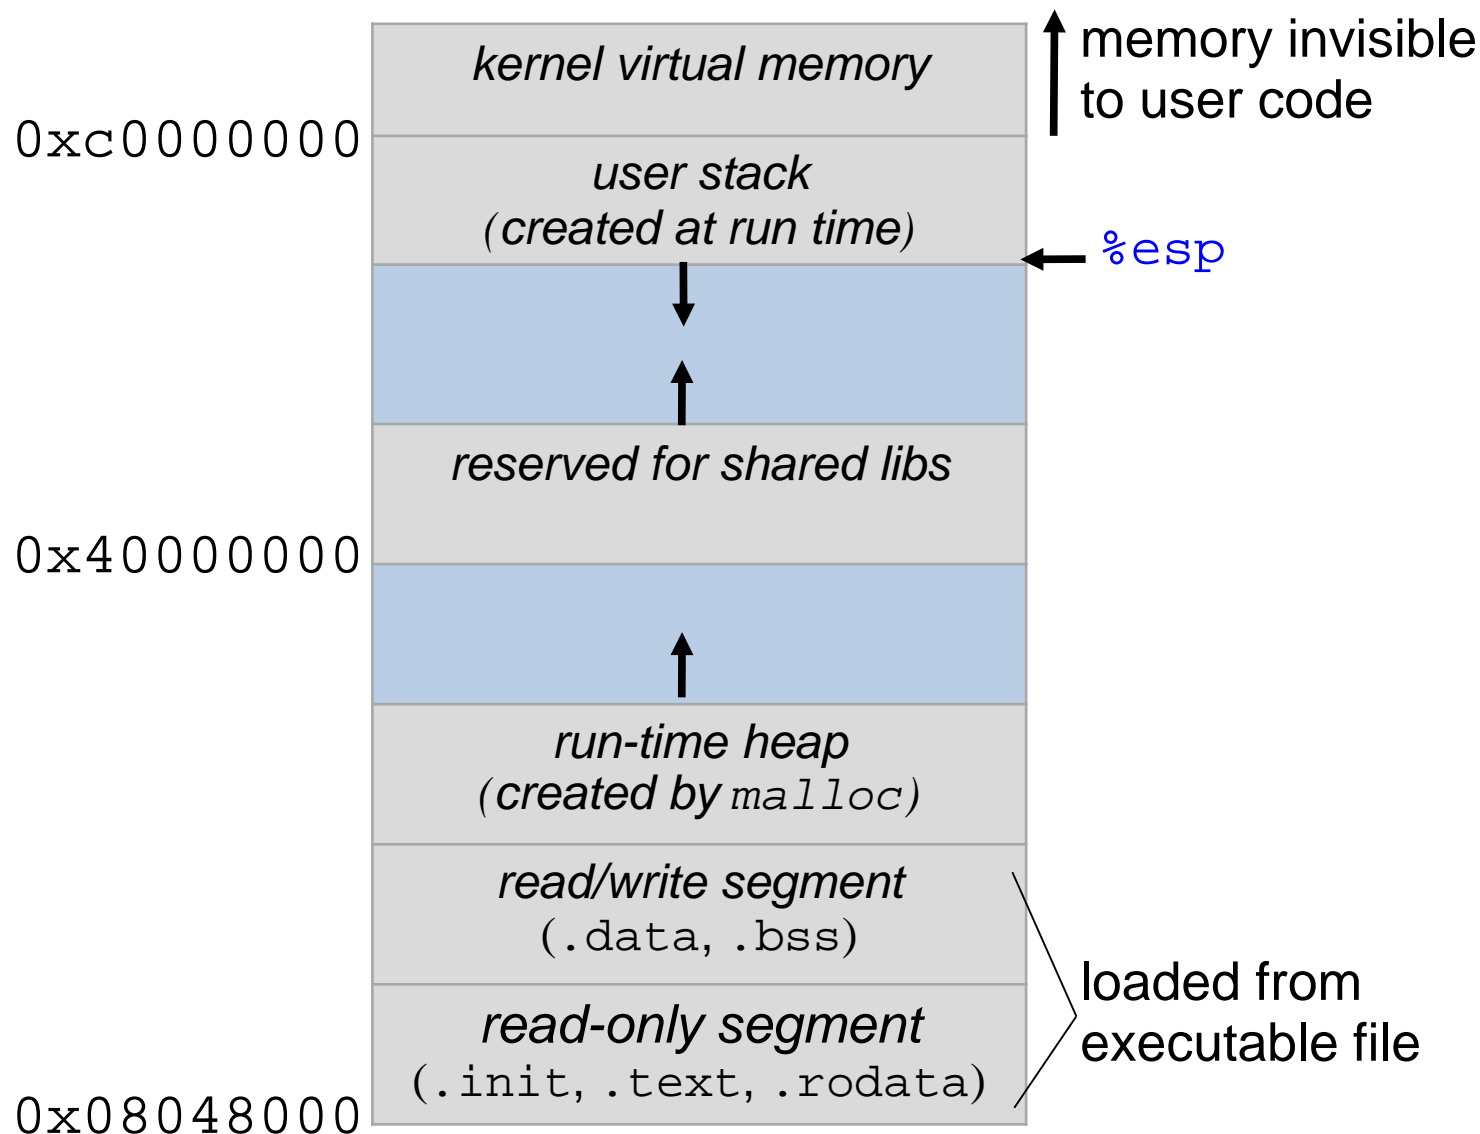
- Because `p` is not a built-in shell command, the shell assumes that `p` is an executable object file.
- The shell invokes loader (by calling function `execve`)
 - copy the code and data from `p` into memory and
 - run the program by jumping to the “entry point”

Loading Executable Object Files

\$./p

- Because `p` is not a built-in shell command, the shell assumes that `p` is an executable object file.
- The shell invokes loader (by calling function `execve`)
 - copy the code and data from `p` into memory and
 - run the program by jumping to the “entry point”
- When the loader runs, it creates a memory image and copies chunks of the executable into the code and data segments.

Unix Run-Time Memory Image



Shared Libraries

- Static libraries sometimes need to be updated
 - Bummer because all programs using the lib need to re-link
- Almost all C programs reference standard I/O functions, and code for these functions appears in the text segment of nearly every running program—waste of memory.

Shared Libraries

- ***Shared library***—an object module that can be loaded and linked with a program in memory, all at load or run time.
 - The process of linking a shared library is called dynamic linking.
- AKA shared objects (.so Unix suffix, DLLs on Microsoft)

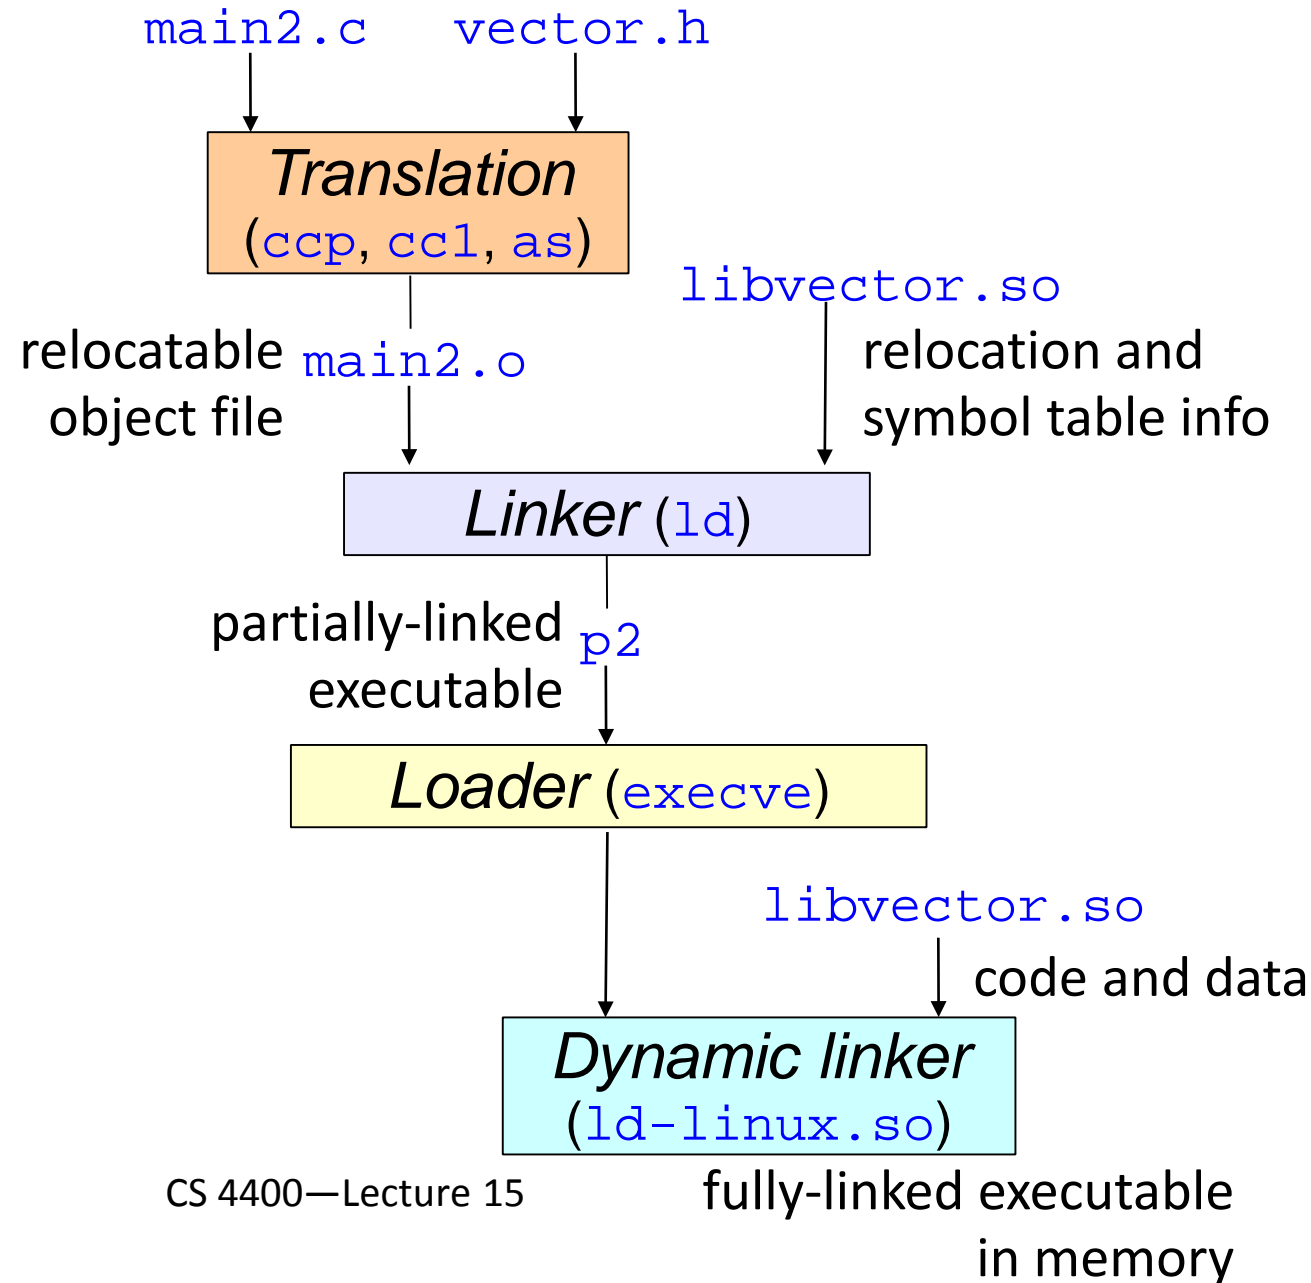
Dynamic Linking

- Shared libraries must be compiled as “position independent code” using the `-fPIC` compiler option
 - PIC uses a dedicated relocation register to figure out where things are
 - Slight loss of performance

```
unix> gcc -o p2 main2.c ./libvector.so
```

- creates `p2` in a form to be linked with `libvector.so` at load time
- Does some of the linking statically and then completes linking process when the program is loaded.

Dynamic Linking w/ Shared Libs



- None of code and data from `libvector.so` is copied into `p2`
- Instead it copies only some relocation and symbol table info to allow references to be resolved at run time.
- Loader notices an `.interp` section with the dynamic linker's path. It passes control there to finish linking. Then passes control to the program.

Run-Time Loading and Linking

- A program requests that the dynamic linker load and link shared libraries while the program is running
 - without having to (partially) link in the libraries at compile time
- Microsoft uses shared libraries to distribute SW updates
 - users download updates and the next time their application runs, it will automatically link and load the new shared library
- Web servers generate dynamic content
 - the appropriate function is loaded/linked at run time
- (See the text for a discussion of the simple interface for the dynamic linker that is provided on Unix systems.)

Summary

- Linkers manipulate object files at compile time (relocatable, static linking), load time or run time (shared libraries, dynamic linking)
- Two main tasks: symbol resolution and relocation
 - each global symbol in an object file is bound to a unique definition
 - the ultimate memory address for each symbol is determined

Summary

- Static linkers combine multiple relocatable object files into a single executable object file at compile time.
- Dynamic linkers are invoked at load or run time to resolve references in user code with definitions in shared libraries.

Summary

- The left-to-right scan of input object files can also be confusing
- The rules linkers use for silently resolving multiple definitions can introduce subtle bugs