

CS 4400

Computer Systems

LECTURE 17

More on process control

Signals

Nonlocal jumps

Question – *Review fork()*

```
#include "csapp.h"

void doit() {
    if(Fork() == 0) {
        Fork();
        printf("hello\n");

        if(Fork() != 0)
            exit(0);
    }
}

int main() {
    doit();
    printf("hello\n");
    exit(0);
}
```

How many “hello”
output lines does this
program print?

Reaping Child Processes

- When a process terminates, the kernel does not remove it from the system immediately.
- The process is retained in a terminated state until it is *reaped* by its parent.
 - a terminated process not yet reaped is called a *zombie*
- If the parent terminates without reaping its children, the kernel arranges for the `init` process to reap them.
 - `init` has PID 1 and is created during system initialization
 - long running programs (i.e., shells) *should always* reap their zombie children because they consume system memory

waitpid Function

- A process waits for its children to terminate by calling

encodes info about child
determines members of the wait set | *modifies default behavior*
`pid_t waitpid(pid_t pid, int* status, int options);`

- By default, `waitpid` suspends execution of the calling process until a child process in its *wait set* terminates.
 - if a process in the wait set has already terminated, `waitpid` returns immediately
 - returns the PID of the terminated child causing `waitpid` to return
 - terminated child is then removed from the system

Determining the Wait Set

- If `pid > 0`, then the wait set is the singleton child process whose PID is equal to `pid`. If `pid = -1`, then the wait set consists of all of the parent's child processes.
- Standard macros interpret the value of `status`.
 - `WIFEXITED(status)` is true if child terminated normally
 - `WEXITSTATUS(status)` returns exit status of child
 - see text for more macros
- If there are no children, `waitpid` returns -1 and `errno` set to `ECHILD`.
 - also returns -1 if interrupted by a signal (`errno` set to `EINTR`)

Example: waitpid

Will the children always be reaped “in order”?

```
/* waitpid1.c */  
  
#include "csapp.h"  
#define N 2
```

```
int main() {  
    int status, i;  
    pid_t pid;
```

```
    for(i = 0; i < N; i++)  
        if((pid = Fork()) == 0) /* child */  
            exit(100+i);
```

```
    /* parent waits for all of its children to terminate */  
    while((pid = waitpid(-1, &status, 0)) > 0) {  
        if(WIFEXITED(status))  
            printf("child %d terminated normally with exit status=%d\n",  
                pid, WEXITSTATUS(status));
```

```
        else  
            printf("child %d terminated abnormally\n", pid);
```

```
    }  
    if(errno != ECHILD)  
        unix_error("waitpid error");
```

```
    exit(0);
```

```
}
```

```
unix> ./waitpid1  
child 22966 terminated normally with exit status=100  
child 22967 terminated normally with exit status=101
```

Question

```
#include "csapp.h"

int main() {
    int status;
    pid_t pid;

    printf("Hello\n");
    pid = Fork();
    printf("%d\n", !pid);
    if(pid != 0)
        if(waitpid(-1, &status, 0) > 0)
            if(WIFEXITED(status) != 0)
                printf("%d\n", WEXITSTATUS(status));

    printf("Bye\n");
    exit(2);
}
```

What does this print? Is output the same every time?

Question

```
#include "csapp.h"

/* Wait() = Waitpid() with pid and options set to
   defaults; it blocks until any child terminates. */

int main() {
    if(Fork() == 0) {
        if(Fork() == 0)
            printf("a");
        else {
            pid_t pid; int status;
            if((pid = Wait(&status)) > 0)
                printf("b");
        }
    }
    else {
        printf("c");
        exit(0);
    }
    printf("d");
    return 0;
}
```

Which outputs are possible?

acdbd

adbdc

abddc

cadbd

bdadc

sleep and pause

- `sleep` suspends a process for some period of time.

```
unsigned int sleep(unsigned int secs);
```

- returns 0 if the requested amount of time has already elapsed
- otherwise, returns number of seconds left to sleep (will happen if it was interrupted by a signal)

Don't try to use this function to ensure that one thing happens before another

- `pause` puts calling function to sleep until a signal is received by the process.

```
int pause(void);
```

Don't use this function in a real program; use `sigsuspend`

execve Function

- Loads and runs a new program in the context of the current process.

executable object file
argument list
environment variable list

```
int execve(char* filename, char* argv[], char* envp);
```

- `execve` returns to calling program only if there's an error.
 - called once, never returns
- `argv` and `envp` each point to a NULL-terminated array of pointers to strings.
 - by convention, `argv[0]` = name of the executable object file
 - each environment variable string has form "NAME=VALUE"

Example: argv and envp

```
/* myecho.c */

#include "csapp.h"

int main(int argc, char* argv[], char* envp[]) {
    int i;

    printf("Command line arguments:\n");
    for(i = 0; i < argc; i++)
        printf("\t argv[%2d]: %s\n", i, argv[i]);

    printf("Environment variables:\n");
    for(i = 0; envp[i] != NULL; i++)
        printf("\t envp[%2d]: %s\n", i, envp[i]);

    exit(0);
}
```

(See text for functions that manipulate `envp`.)

```
[user@lab1-3 ~]$ ./a.out hi
Command line arguments:
    argv[ 0]: ./a.out
    argv[ 1]: hi
Environment variables:
    envp[ 0]: HOSTNAME=lab1-3
    envp[ 1]: MALLOC_CHECK_ =1
    envp[ 2]: TERM=xterm-256color
    envp[ 3]: SHELL=/bin/bash
    envp[ 4]: HISTSIZE=1000
    ...
```

Programs vs. Processes

- *Program*—collection of code and data
- *Process*—a specific instance of a program in execution
- `fork` runs the same program in a new child process that is a duplicate of the parent process.
- `execve` loads and runs a new program in context of the current process and *does not create a new process*.
 - new program has same PID
 - inherits all of the file descriptors that were open at the time of the call to `execve`

`execve`, to perform a sequence of read/evaluate steps.

Shells

- Read step—read a command line from the user.
- Evaluate step—parse the command line and run programs on the behalf of the user.
- *Simple shell example:*

```
int main() {
    char cmdline[MAXLINE];

    while(1) {
        printf("> ");
        Fgets(cmdline, MAXLINE, stdin);
        if (feof(stdin))
            exit(0);

        eval(cmdline);
    }
}
```

```

int parseline(char* buf, char** argv);
int builtin_command(char** argv);

void eval(char *cmdline) {      /* evaluate a command line */
    char *argv[MAXARGS]; /* argv for execve() */
    char buf[MAXLINE];      /* holds modified command line */
    int bg;                  /* should the job run in bg or fg? */
    pid_t pid;               /* process id */

    strcpy(buf, cmdline);
    bg = parseline(buf, argv); /* true if last argv is & */
    if(argv[0] == NULL) return; /* ignore empty lines */

    if(!builtin_command(argv)) {
        if((pid = Fork()) == 0) /* child runs user job */
            if(execve(argv[0], argv, environ) < 0) {
                printf("%s: Command not found.\n", argv[0]);
                exit(0);
            }

        /* parent waits for foreground job to terminate */
        if(!bg) {
            int status;
            if(waitpid(pid, &status, 0) < 0)
                unix_error("waitfg: waitpid error");
        }
        else
            printf("%d %s", pid, cmdline);
    }
    return;
} /* shell is flawed because children not reaped */

```

Signals

- *Signal*—a message that notifies a process that an event of some type has occurred in the system.
 - allows processes to interrupt other processes
- Transfer of a signal to a destination process:
 1. Kernel *sends* a signal to a destination process by updating some state in the context of the destination process.
 2. A destination process *receives* a signal when it is forced by the kernel to react (ignore signal, terminate, or catch signal) to the delivery of the signal.
- (See text for a list of Linux signals.)

Pending Signals

- *Pending signal*—sent but not yet received.
- At any point, there can be at most one pending signal of a particular type.
- If a process p has a pending signal of type k , any subsequent signals of type k sent to p are discarded.
- A process can selectively block receipt of certain signals (signal is delivered, but not received until unblocked).
- A pending signal is received at most once.
- Kernel keeps track of pending and blocked signals.

Process Groups

- Every process belongs to exactly one *process group*.
 - a process group is identified by a process group ID > 0
 - `pid_t getpgrp(void)` returns process group ID of current process
- By default, a child process belongs to the process group of its parent.
- `setpgid` changes the process group of `pid` to `pgid`.

```
pid_t setpgid(pid_t pid, pid_t pgid);
```

- if `pid=0`, PID of current process is used
- if `pgid=0`, PID of process specified by `pid` is used for group id
- what does `setpgid(0, 0)` do?

Sending Signals

- `kill` sends signal number `sig` to other process(es).
- `int kill(pid_t pid, int sig);`
 - if `pid > 0`, sends to process `pid`
 - if `pid < 0`, sends to every process in process group `abs(pid)`

```
#include "csapp.h"

int main() {
    pid_t pid;

    /* child sleeps until SIGKILL signal received
       then dies */
    if((pid = Fork()) == 0) {
        Pause(); /* wait for signal */
        printf("control never reaches here");
        exit(0);
    }

    /* parent sends SIGKILL signal to child */
    Kill(pid, SIGKILL);
    exit(0);
}
```

Receiving Signals

- When the kernel is ready to pass control to process p , it checks the set of pending, unblocked signals.
 - if the set is empty, continue with l_{next} in p
 - otherwise, choose some signal number k (usually the smallest) from the set and force p to receive the signal
- The process completes some *action* in response and then control passes to l_{next} .
- Each signal has a default action (see text). Process either terminates, terminates and dumps core, stops until restarted by `SIGCONT` signal, or ignores signal.

Modifying Default

- `signal` modifies the default action for a signal.
 - `handler_t* signal(int signum, handler_t* handler);`
 - `handler` is the address of a user-defined function
 - default actions of `SIGSTOP` and `SIGKILL` cannot be changed
 - generally, `sigaction()` should be used instead of `signal()`

```
#include "csapp.h"

void handler(int sig) { /* SIGINT handler */
    printf("Caught SIGINT\n");
    exit(0);
}

int main() {
    /* Install SIGINT handler */
    if(signal(SIGINT, handler) == SIG_ERR)
        unix_error("signal error");

    pause(); /* Wait for ctrl-c from keyboard */

    exit(0);
}
```

Explicitly Blocking Signals

- `sigprocmask` explicitly blocks selected signals.

```
int sigprocmask(int how, sigset_t* set, sigset_t* oldset);
```

- The set of blocked signals is maintained as a bit vector
`blocked`.

- Behavior depends on argument `how`.

- `SIG_BLOCK`—adds signals in `set` to `blocked`

`(blocked |= set)`

- `SIG_UNBLOCK`—removes signals in `set` from `blocked`

`(blocked &= ~set)`

- `SIG_SETMASK`—`blocked = set`

```

void handler(int sig) {
    pid_t pid;
    while((pid = waitpid(-1, NULL, 0)) > 0) /* Reap a zombie child */
        deletejob(pid); /* Delete the child from the job list */
    if(errno != ECHILD)
        unix_error("waitpid error");
}

int main(int argc, char** argv) {
    int pid;
    sigset_t mask;

    Signal(SIGCHLD, handler);
    initjobs(); /* Initialize job list (to keep track of children) */

    while(1) {
        Sigemptyset(&mask);
        Sigaddset(&mask, SIGCHLD);
        Sigprocmask(SIG_BLOCK, &mask, NULL); /* Block SIGCHLD */

        /* Child process */
        if((pid = Fork()) == 0) {
            Sigprocmask(SIG_UNBLOCK, &mask, NULL); /* Unblock SIGCHLD */
            Execve("/bin/ls", argv, NULL);
        }

        /* Parent process */
        addjob(pid); /* Add the child to the job list */
        Sigprocmask(SIG_UNBLOCK, &mask, NULL); /* Unblock SIGCHLD */
    }
    exit(0);
}

```

Nonlocal Jumps

- Transfer control from one function to another currently executing function, without having to go through the normal call-and-return sequence.

- `setjmp` saves the current stack context in `env`.

```
int setjmp(jmp_buf env);
```

- `longjmp` restores the stack context from the `env` buffer and then triggers a return from the most recent `setjmp` call that initialized `env`.

```
int longjmp(jmp_buf env, int retval);
```

- `setjmp` then returns with return value `retval`

Nonlocal Jumps

- `setjmp` is called once and returns multiple times.
 - once when it is first called and stack context is saved
 - once for each corresponding call to `longjmp`
- `longjmp` never returns.
- Nonlocal jumps permit
 - immediate return from a deeply-nested function call, usually as a result of detecting some error (return directly to an error handler, rather than unwinding the call stack)
 - branching out of a signal handler to a specific code location, rather than returning to the instruction that was interrupted at the arrival of the signal


```
jmp_buf buf;

int error1 = 0;
int error2 = 1;

void foo(void), bar(void);

int main() {
    int rc;

    rc = setjmp(buf); /* returns 0 when called directly */
    if(rc == 0)        /* returns !=0 when called indirectly */
        foo();
    else if(rc == 1)
        printf("Detected an error1 condition in foo\n");
    else if(rc == 2)
        printf("Detected an error2 condition in foo\n");
    else
        printf("Unknown error condition in foo\n");
    exit(0);
}

void foo(void) { /* deeply nested function foo */
    if(error1)
        longjmp(buf, 1);
    bar();
}

void bar(void) {
    if(error2)
        longjmp(buf, 2);
}
```

```

/* restart.c */

sigjmp_buf buf;

void handler(int sig) {
    siglongjmp(buf, 1); /* version of longjmp that can be */
}                       /* used by signal handlers */
                        /* 1 means to restore the signal mask */

int main() {
    Signal(SIGINT, handler);

    if(!sigsetjmp(buf, 1)) /* version of setjmp for sig handlers */
        printf("starting\n"); /* 1 means to save the signal mask */
    else
        printf("restarting\n");

    while(1) {
        Sleep(1);
        printf("processing...\n");
    }
    exit(0);
}

```

```

unix> ./restart
starting
processing...
processing...
restarting      user types ctrl-c
processing...
restarting      user types ctrl-c
processing...

```

Summary

- ECF occurs at all levels of a computer system.
- *Hardware level*: interrupt, trap, fault, and abort classes of exceptions.
- *OS level*: a process provides the illusion that a program has exclusive use of the processor and memory.
- *Application level*: apps can create and wait for child processes, run new programs, and catch signals from other processes.
 - C programs can use nonlocal jumps to bypass the normal call/return stack discipline and branch directly to a function.