

CS 4400: Computer Systems

Fall 2014

Midterm Exam 1

- Please give your solutions in the space provided on the exam. If you choose to show your work on the exam, be sure to clearly indicate your solution to each problem.
- The exam is closed book, but you are allowed a 1 page cheat sheet, no electronics.
- The number of points assigned to each problem is indicated at the top of the page.
- Unless a problem indicates otherwise, you should assume that C code will be compiled on an IA32 platform where, for example, an `int` is 32 bits
- Make sure that you have 7 numbered pages including this cover sheet.

Name: _____

uUID: _____

Given the following variable definitions, list the type and value for the following statements.

```
int      i = -1;
unsigned u = 20;
float    f = 1.0e-02;
double   d = 2.0e20;
```

Expression	Type	Value
<code>i >> 2</code>	<code>int</code>	<code>-1</code>
<code>u << 4</code>	<code>unsigned</code>	<code>320</code>
<code>i + u</code>	<code>unsigned</code>	<code>19</code>
<code>f + i</code>	<code>float</code>	<code>-0.99</code>
<code>d + f</code>	<code>double</code>	<code>2.0e20</code>
<code>d * f</code>	<code>double</code>	<code>2.0e18</code>

Consider a 7-bit two's complement representation. Fill in the empty boxes in the following table. You need not fill in entries marked —. For this problem you are doing 2's complement operations, not evaluating C code, so there is no undefined behavior. Evaluate the results of all expressions using 2's complement arithmetic. Be aware of overflow behavior.

$$2^7 = 128, 2^6 = 64, 2^5 = 32, 2^4 = 16, 2^3 = 8, 2^2 = 4, 2^1 = 2, 2^0 = 1$$

Expression	Decimal Representation	7-bit Binary Representation
—	19	0010011
—	-54	1001010
TMin	-64	1000000
TMax	63	0111111
—	37	0100101
—	-36	1011100
TMax-TMin	-1	1111111
TMin-TMax	1	0000001
TMin + 1	-63	1000001
-TMin	-64	1000000

Consider a 7-bit two's complement representation. Fill in the empty boxes in the following table. You need not fill in entries marked —. For this problem you are doing 2's complement operations, not evaluating C code, so there is no undefined behavior. Evaluate the results of all expressions using 2's complement arithmetic. Be aware of overflow behavior.

$$2^7 = 128, 2^6 = 64, 2^5 = 32, 2^4 = 16, 2^3 = 8, 2^2 = 4, 2^1 = 2, 2^0 = 1$$

Expression	Decimal Representation	7-bit Binary Representation
—	47	0101111
—	-9	1110111
TMin	-64	1000000
TMax	63	0111111
-TMin	-64	1000000
—	55	0110111
—	-6	1111010
TMin + 1	-63	1000001
TMax-TMin	-1	1111111
TMin-TMax	1	0000001

Consider a 7-bit two's complement representation. Fill in the empty boxes in the following table. You need not fill in entries marked —. For this problem you are doing 2's complement operations, not evaluating C code, so there is no undefined behavior. Evaluate the results of all expressions using 2's complement arithmetic. Be aware of overflow behavior.

$$2^7 = 128, 2^6 = 64, 2^5 = 32, 2^4 = 16, 2^3 = 8, 2^2 = 4, 2^1 = 2, 2^0 = 1$$

Expression	Decimal Representation	7-bit Binary Representation
TMin	-64	1000000
TMax	63	0111111
—	30	0011110
—	-41	1010111
TMin + 1	-63	1000001
-TMin	-64	1000000
—	37	0100101
—	-51	1001101
TMax-TMin	-1	1111111
TMin-TMax	1	0000001

Consider the following 13-bit floating-point number based on IEEE format.

- The most significant bit indicates the sign.
- The next 4 bits are the exponent.
- The last 8 bits are the fraction.
- The representation encodes numbers of the form: $V = (-1)^s \times M \times 2^E$, where M is the significand and E is the biased exponent.

Fill in the table below. The following are the instructions for each field.

- Hex: The 13-bit binary representation, given in 4-digit hexadecimal.
- M: The value of the significand. This should be fractional form number.
- E: The integer value of the exponent.
- Value: The numeric value represented.

Note: You need not fill in entries marked with –.

Description	Hex	M	E	Value
Negative zero	0x1000	0	-6	-0.0
Positive infinity	0x0f00	–	–	inf
-1.125	0x1720	$1 + \frac{1}{8}$	0	-1.125
NaN	0x1f10	–	–	NaN
2.0	0x0800	1	1	2.0
Smallest denormalized > 0	0x0001	$\frac{1}{256}$	-6	0.000061
Largest normalized > 0	0x0eff	$1 + \frac{255}{256}$	7	255.5

Consider the following 13-bit floating-point number based on IEEE format.

- The most significant bit indicates the sign.
- The next 4 bits are the exponent.
- The last 8 bits are the fraction.
- The representation encodes numbers of the form: $V = (-1)^s \times M \times 2^E$, where M is the significand and E is the biased exponent.

Fill in the table below. The following are the instructions for each field.

- Hex: The 13-bit binary representation, given in 4-digit hexadecimal.
- M: The value of the significand. This should be fractional form number.
- E: The integer value of the exponent.
- Value: The numeric value represented.

Note: You need not fill in entries marked with –.

Description	Hex	M	E	Value
Negative zero	0x1000	0	-6	-0.0
Positive infinity	0x0f00	–	–	inf
2.5	0x0840	$1 + \frac{1}{4}$	1	2.5
NaN	0x1f10	–	–	NaN
2.0	0x0800	1	1	2.0
Smallest denormalized > 0	0x0001	$\frac{1}{256}$	-6	0.000061
Largest normalized > 0	0x0eff	$1 + \frac{255}{256}$	7	255.5

Consider the following 13-bit floating-point number based on IEEE format.

- The most significant bit indicates the sign.
- The next 4 bits are the exponent.
- The last 8 bits are the fraction.
- The representation encodes numbers of the form: $V = (-1)^s \times M \times 2^E$, where M is the significand and E is the biased exponent.

Fill in the table below. The following are the instructions for each field.

- Hex: The 13-bit binary representation, given in 4-digit hexadecimal.
- M: The value of the significand. This should be fractional form number.
- E: The integer value of the exponent.
- Value: The numeric value represented.

Note: You need not fill in entries marked with –.

Description	Hex	M	E	Value
Negative zero	0x1000	0	-6	-0.0
Positive infinity	0x0f00	–	–	inf
2.5	0x0840	$1 + \frac{1}{4}$	1	2.5
-1.125	0x1720	$1 + \frac{1}{8}$	0	-1.125
NaN	0x1f10	–	–	NaN
Smallest denormalized > 0	0x0001	$\frac{1}{256}$	-6	0.000061
Largest normalized > 0	0x0eff	$1 + \frac{255}{256}$	7	255.5

Match each of `z ()` functions with one of a, b, c, d, and e on the right. Remember that IA32 return values are passed in `%eax`. Note: `cmov` is a conditional move.

a

```
int z (int x, int y)
{
    if (x>y) return x;
    return y;
}
```

b

```
int z (int x, int y)
{
    if (x+y) return 0;
    return y;
}
```

c

```
int z (int x, int y)
{
    if (x||y) return x;
    return 0;
}
```

d

```
int z (int x, int y)
{
    if (x+1) y++;
    return x-y;
}
```

e

```
int z (int x, int y)
{
    if (!x&&!y) x++;
    return y-x;
}
```

```
a:  pushl    %ebp
      movl    %esp, %ebp
      movl    8(%ebp), %eax
      movl    12(%ebp), %edx
      popl    %ebp
      cmpl    %eax, %edx
      cmovge  %edx, %eax
      ret
```

```
b:  pushl    %ebp
      movl    %esp, %ebp
      movl    12(%ebp), %eax
      movl    8(%ebp), %edx
      popl    %ebp
      addl    %eax, %edx
      testl   %edx, %edx
      movl    $0, %edx
      cmovne  %edx, %eax
      ret
```

```
c:  pushl    %ebp
      movl    %esp, %ebp
      movl    8(%ebp), %eax
      popl    %ebp
      ret
```

```
d:  pushl    %ebp
      xorl    %edx, %edx
      movl    %esp, %ebp
      movl    8(%ebp), %eax
      cmpl    $-1, %eax
      setne   %dl
      addl    12(%ebp), %edx
      popl    %ebp
      subl    %edx, %eax
      ret
```

```
e:  pushl    %ebp
      movl    %esp, %ebp
      movl    12(%ebp), %eax
      movl    8(%ebp), %edx
      testl   %eax, %eax
      jne     .L2
      testl   %edx, %edx
      movl    $1, %ecx
      cmovle  %ecx, %edx
.L2: subl    %edx, %eax
      popl    %ebp
      ret
```

Consider the following x86 assembly code for a function `loopy()`:

```
.globl loopy
.type    loopy, @function
loopy:
    pushl    %ebp
    movl     $1, %ecx
    movl     %esp, %ebp
    movl     $1, %eax
    pushl    %ebx
    movl     8(%ebp), %edx
    movl     12(%ebp), %ebx
    cmpl     %ebx, %edx
    jae .L3
.L6:
    addl     $1, %edx
    addl     %ecx, %ecx
    cmpl     %edx, %ebx
    ja .L6
    movl     %ecx, %eax
.L3:
    popl     %ebx
    popl     %ebp
    ret
```

Based on the assembly code above, fill in the blanks below in its corresponding C source code. You may not refer to registers in the source code.

```
unsigned loopy (unsigned x, unsigned y)
{
    int result = 1;
    while (x < y) {
        result = result*2;
        x = x+1;
    }
    return result;
}
```

Consider the following x86 assembly code for a function `loopy()`:

```
.globl loopy
.type    loopy, @function
loopy:
    pushl    %ebp
    movl     $2, %eax
    movl     %esp, %ebp
    movl     8(%ebp), %edx
    movl     12(%ebp), %ecx
.L2:
    subl     $1, %edx
    sall     $2, %eax
    cmpl     %ecx, %edx
    jb .L2
    popl     %ebp
    ret
```

Based on the assembly code above, fill in the blanks below in its corresponding C source code. You may not refer to registers in the source code.

```
unsigned loopy (unsigned x, unsigned y)
{
    int result = 2;
    do {
        result = result*4;
        x = x-1;
    } while (x < y);
    return result;
}
```

Consider the following x86 assembly code for a function `loopy()`:

```
.globl loopy
.type    loopy, @function
loopy:
    pushl    %ebp
    xorl     %eax, %eax
    movl     %esp, %ebp
    movl     8(%ebp), %edx
    movl     12(%ebp), %ecx
    cmpl     %ecx, %edx
    jae .L3
    subl     %edx, %ecx
    xorl     %edx, %edx
.L4:
    addl     %edx, %eax
    addl     $1, %edx
    cmpl     %ecx, %edx
    jne .L4
.L3:
    popl     %ebp
    ret
```

Based on the assembly code above, fill in the blanks below in its corresponding C source code. You may not refer to registers in the source code.

```
unsigned loopy (unsigned x, unsigned y)
{
    int i;
    int result = 0;
    for (i = 0; x < y; i++) {
        result = result+i;
        x = x+1;
    }
    return result;
}
```